

The Problem

- Response Latency
- Network traffic
- Types of Caches
 - Browser
 - Proxy
 - Gateway

Caching Principle

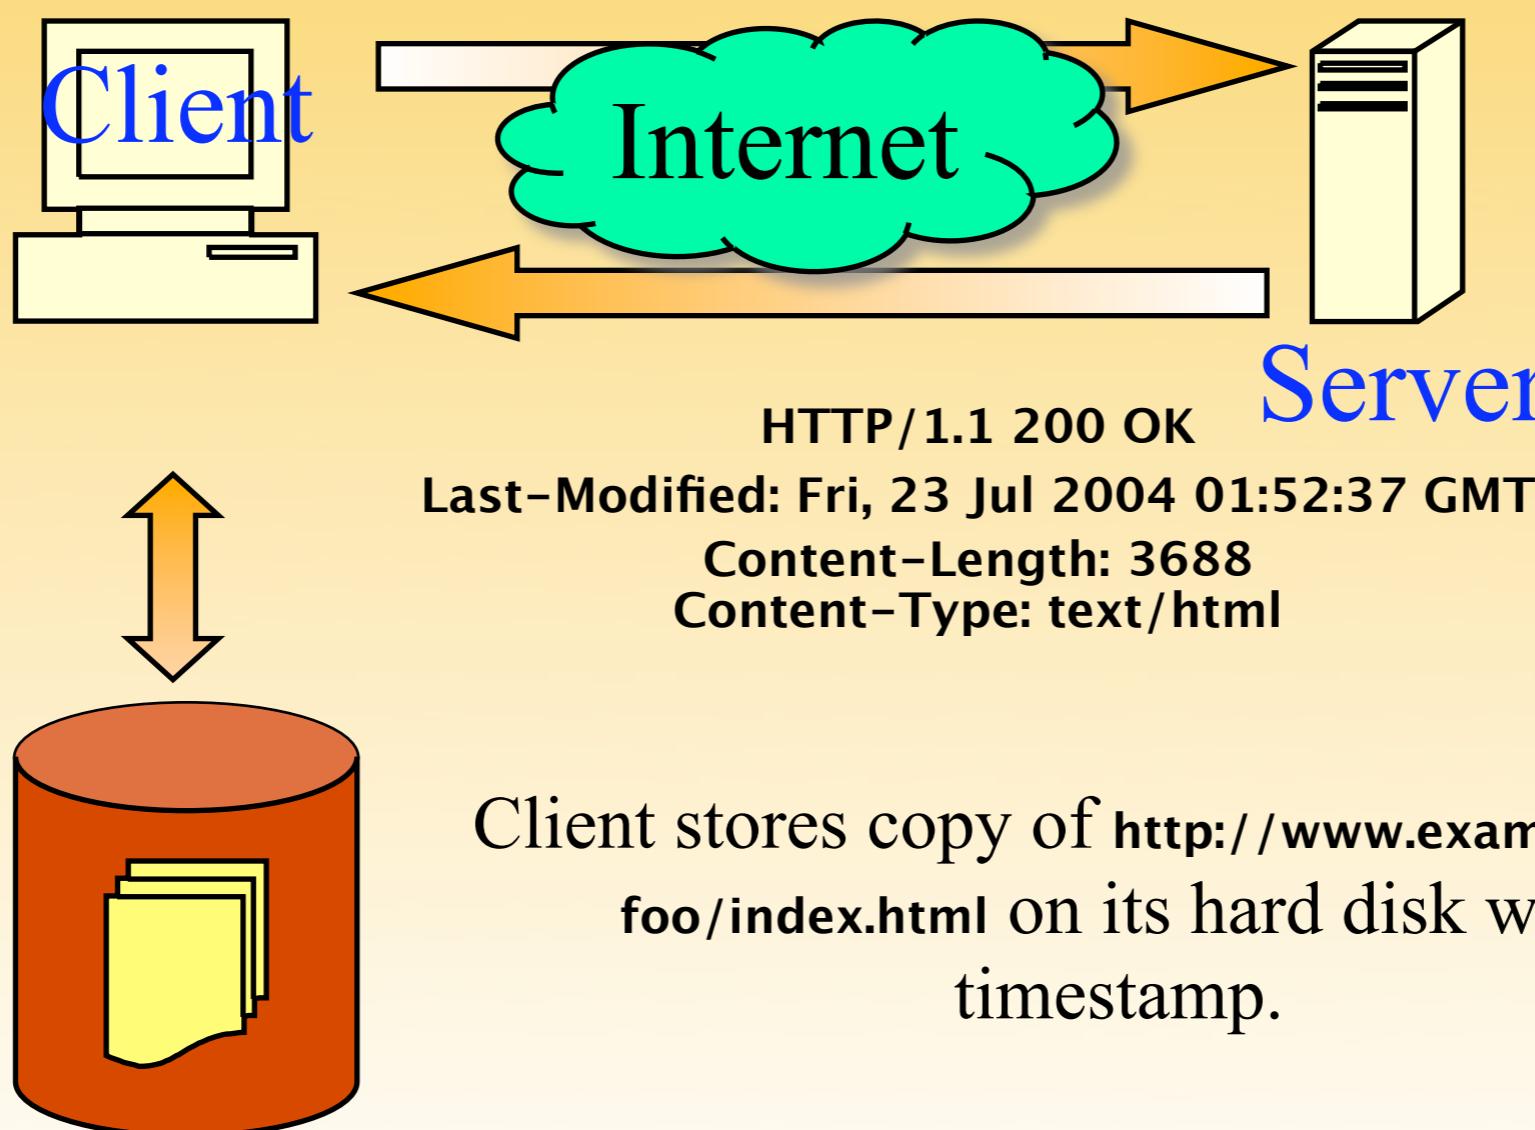
- Fast resources are scarce - e.g. storage close to clients (low propagation delay & fast links)
- Aim to locate commonly used objects in fast resource, other objects can remain in slower resources.
- Determining which objects are commonly accessed:
- Accesses are often correlated, and are said to exhibit “locality”.
 - **Temporal locality:** Information accessed at one time is likely to be accessed again in the near future.
 - **Spatial locality:** Information accessed at one point is likely to be accessed also by adjacent points.

Network Content Caching

- Temporal locality: Content used once is likely to be used again, in the near future, by that node.
 - e.g. view Google logo one day; likely to view again the next day.
- Spatial locality: Content used by one user is likely to also be used by nearby users.
 - e.g. one student @ UM accesses slashdot.org => likely that others will also
- “In one study spanning more than a month, out of all the objects requested by individual users, on average close to 60 percent of those objects were requested more than once by the same user. . . . of the hits recorded in another caching study, up to 85 percent were the result of multiple users requesting the same object.”

Browsers use private caches

GET /foo/index.html HTTP/1.1
Host: www.example.com



Local Cache Control

General Network Update Encryption

Connection

Configure how Firefox connects to the Internet [Settings...](#)

Offline Storage

Use up to MB of space for the cache [Clear Now](#)

Tell me when a website asks to store data for offline use [Exceptions...](#)

The following websites have stored data for offline use:

-

Local Proxy Control

TCP/IP DNS WINS AppleTalk 802.1X Proxies Ethernet

Configure Proxies: Manually

Select a protocol to configure:

- FTP Proxy
- Web Proxy (HTTP)
- Secure Web Proxy (HTTPS)
- Streaming Proxy (RTSP)
- SOCKS Proxy
- Gopher Proxy

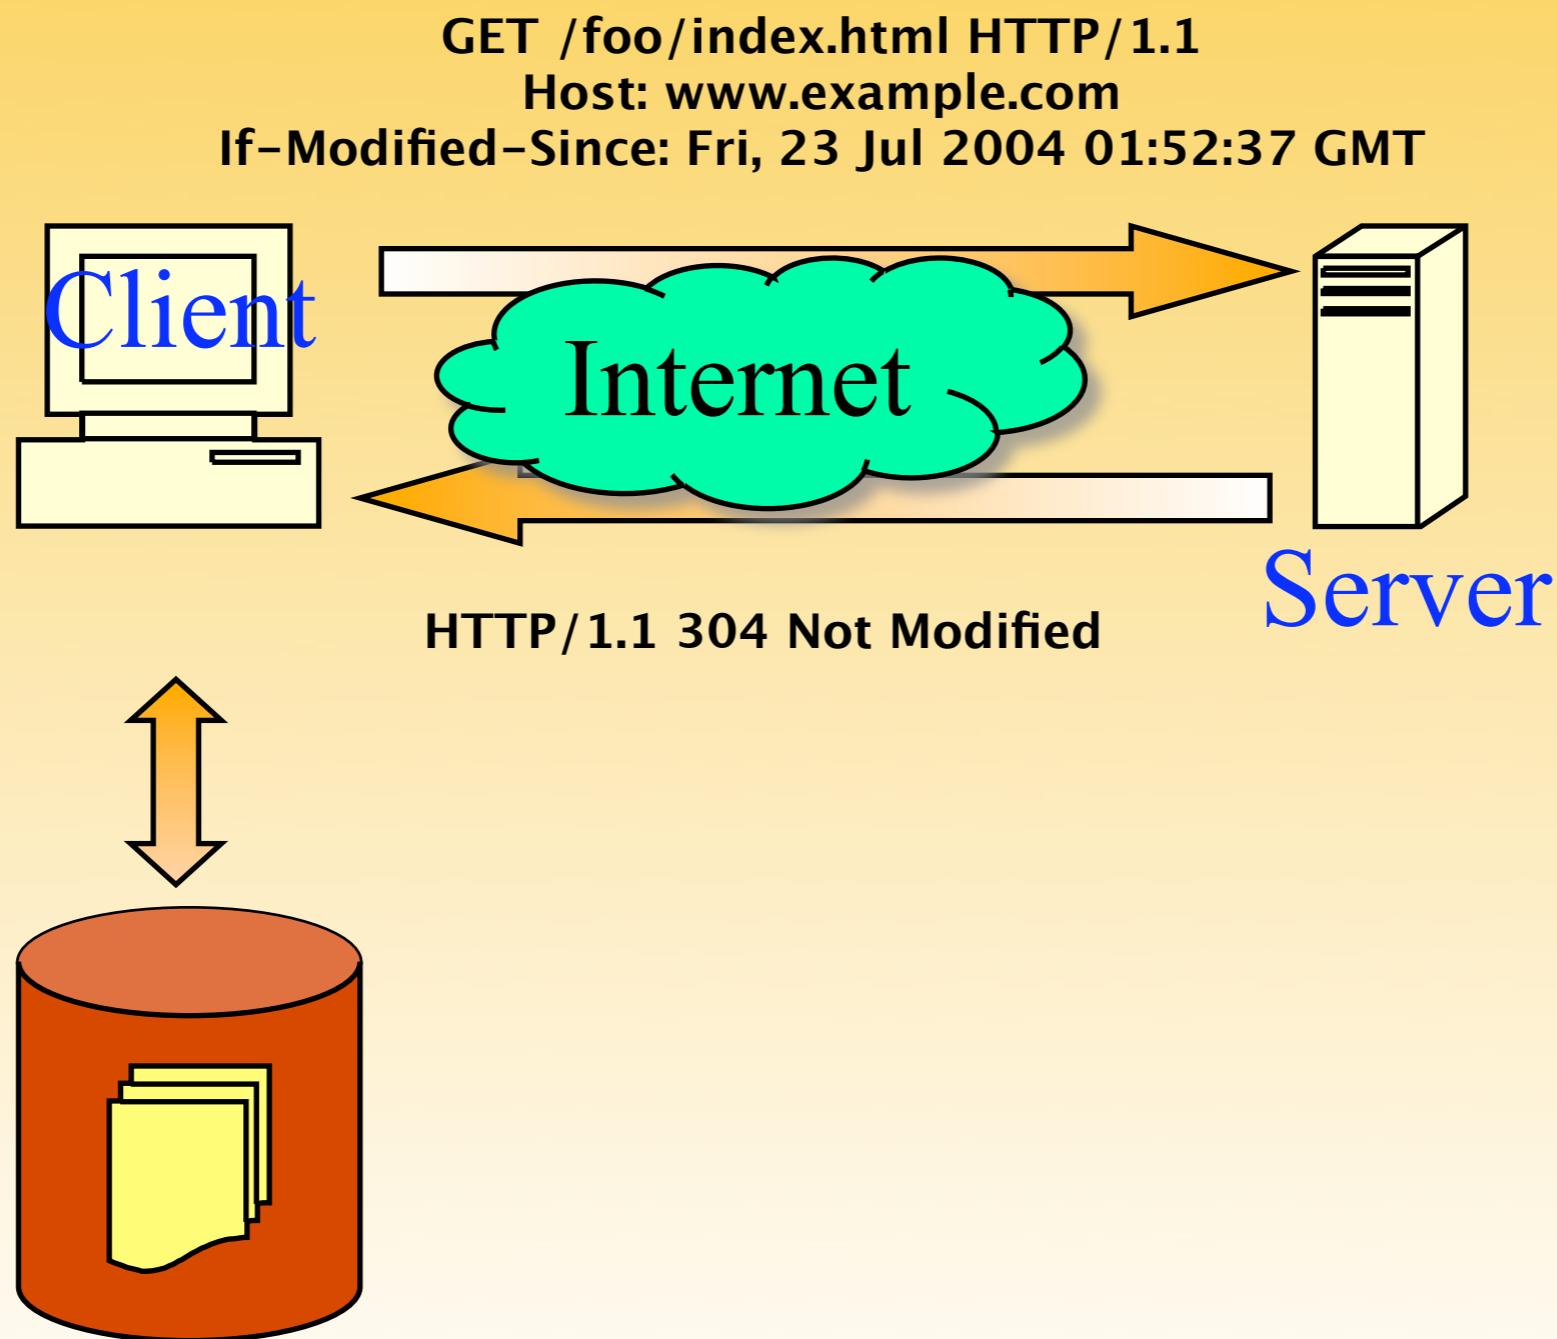
Exclude simple hostnames

Bypass proxy settings for these Hosts & Domains:

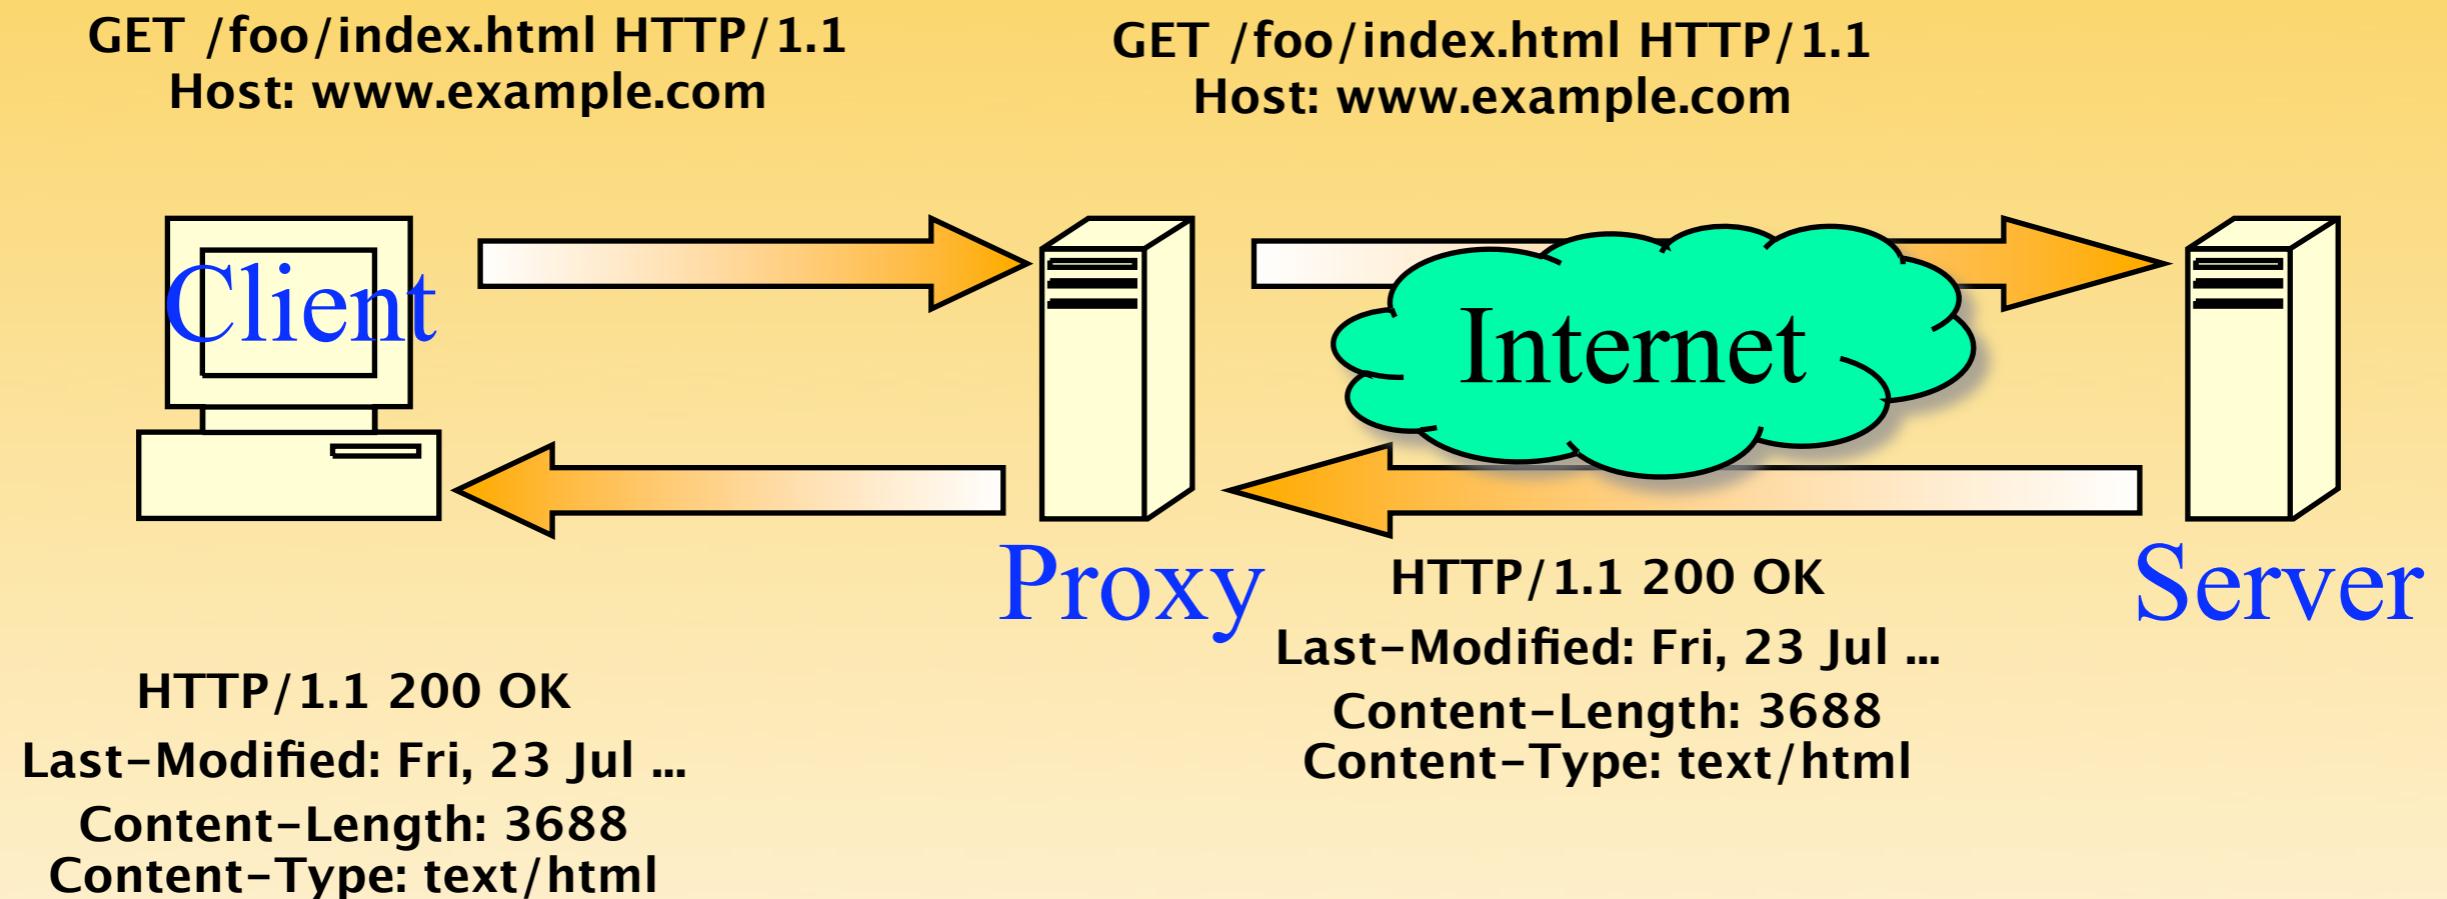
*.local, 169.254/16

Use Passive FTP Mode (PASV)

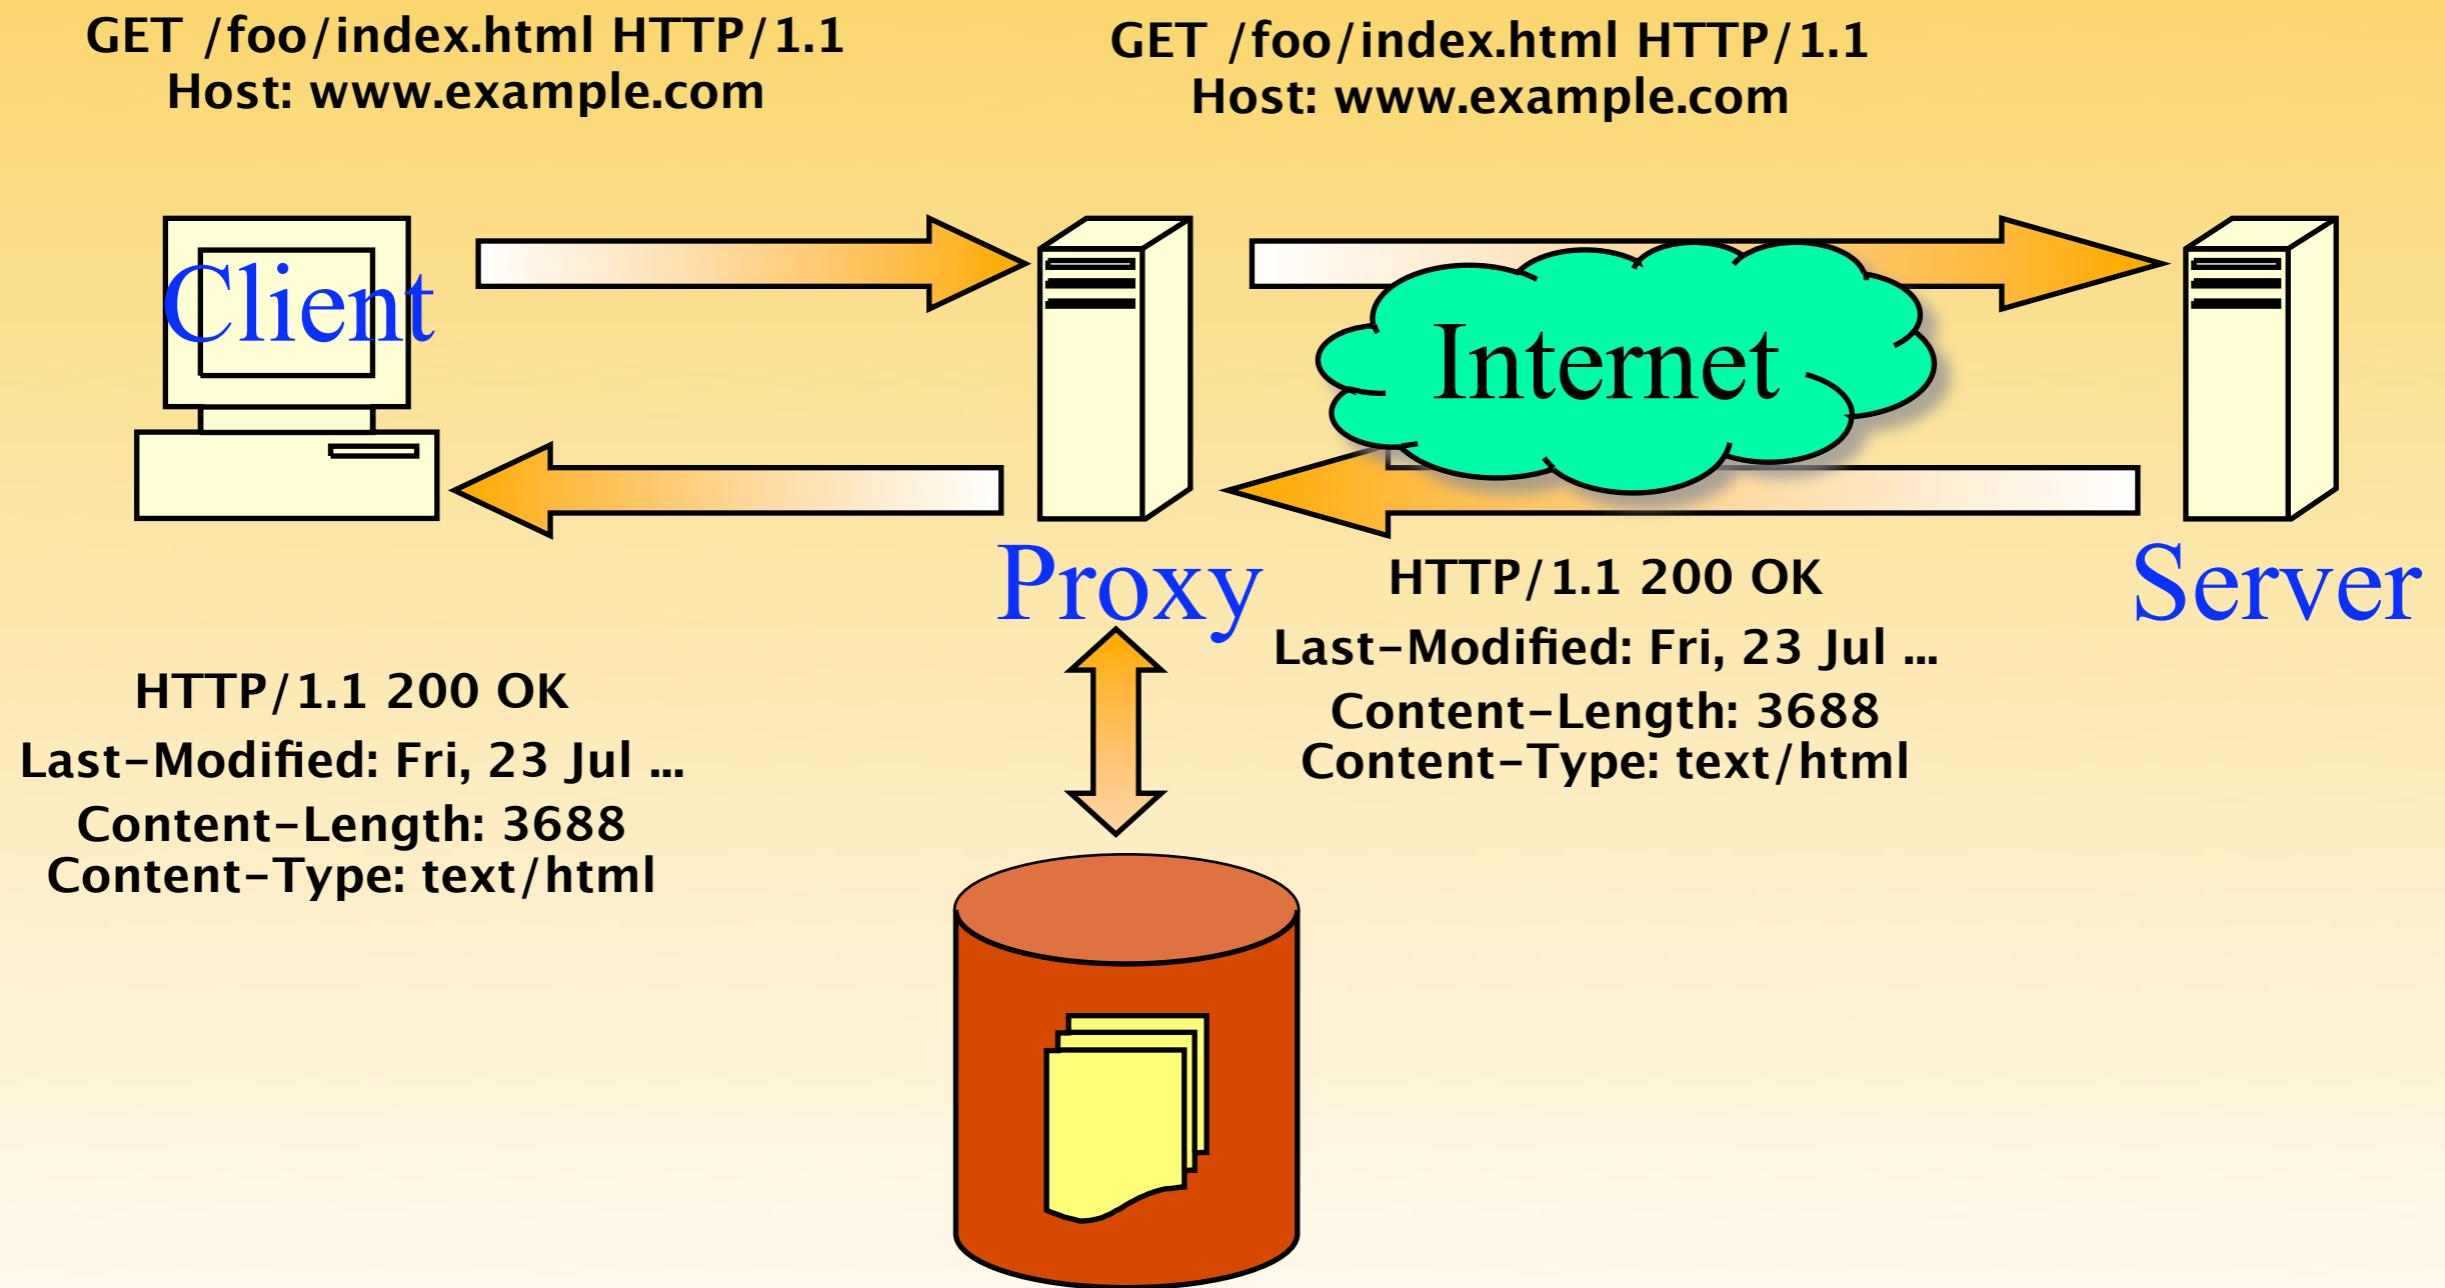
Revalidation (Conditional GET)



Non-Caching Proxy

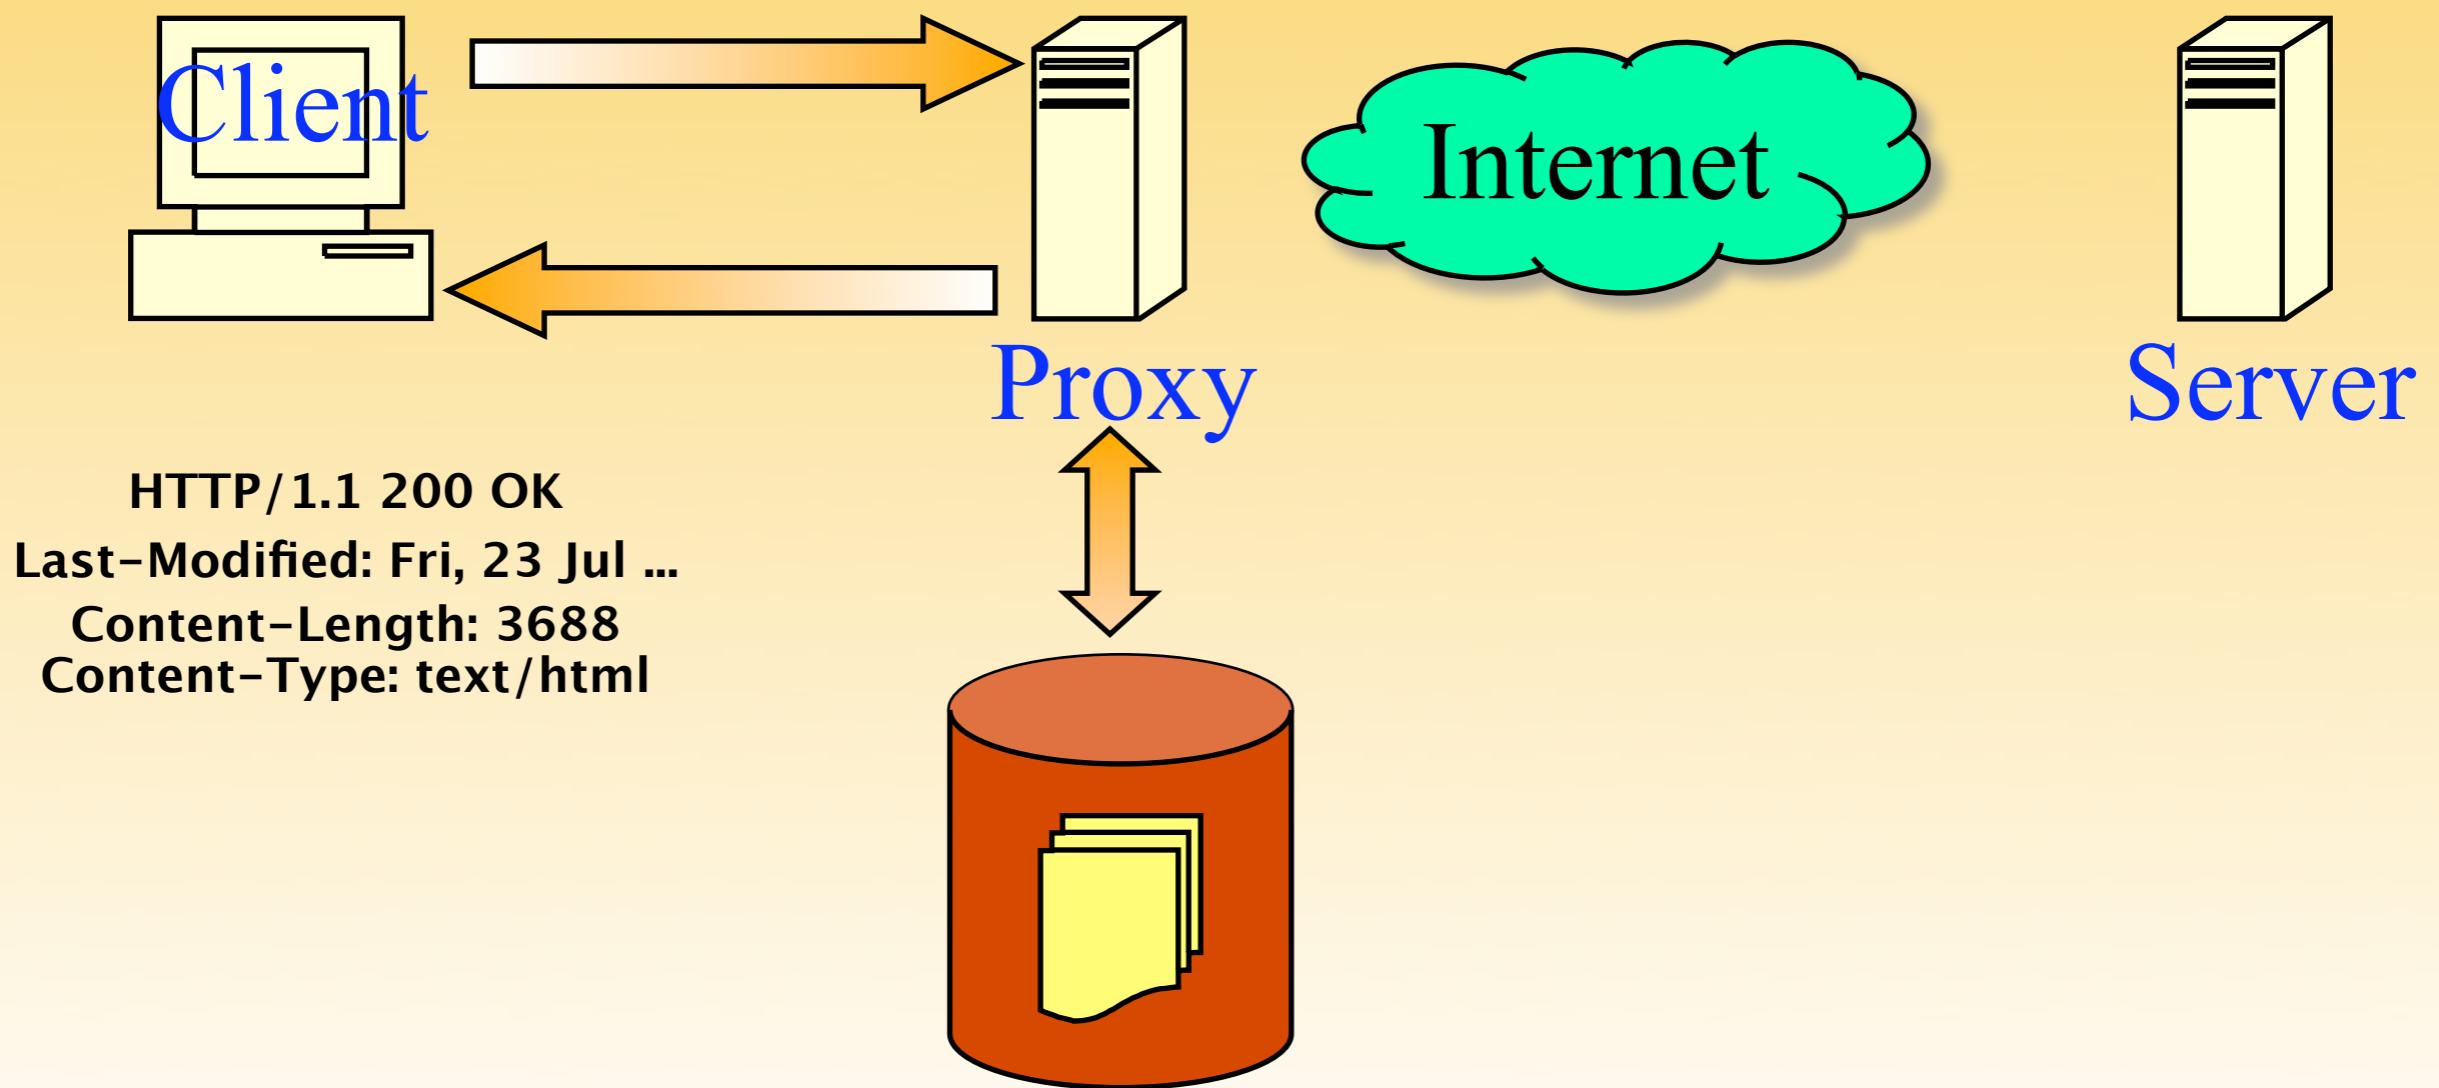


Proxy Cache Miss

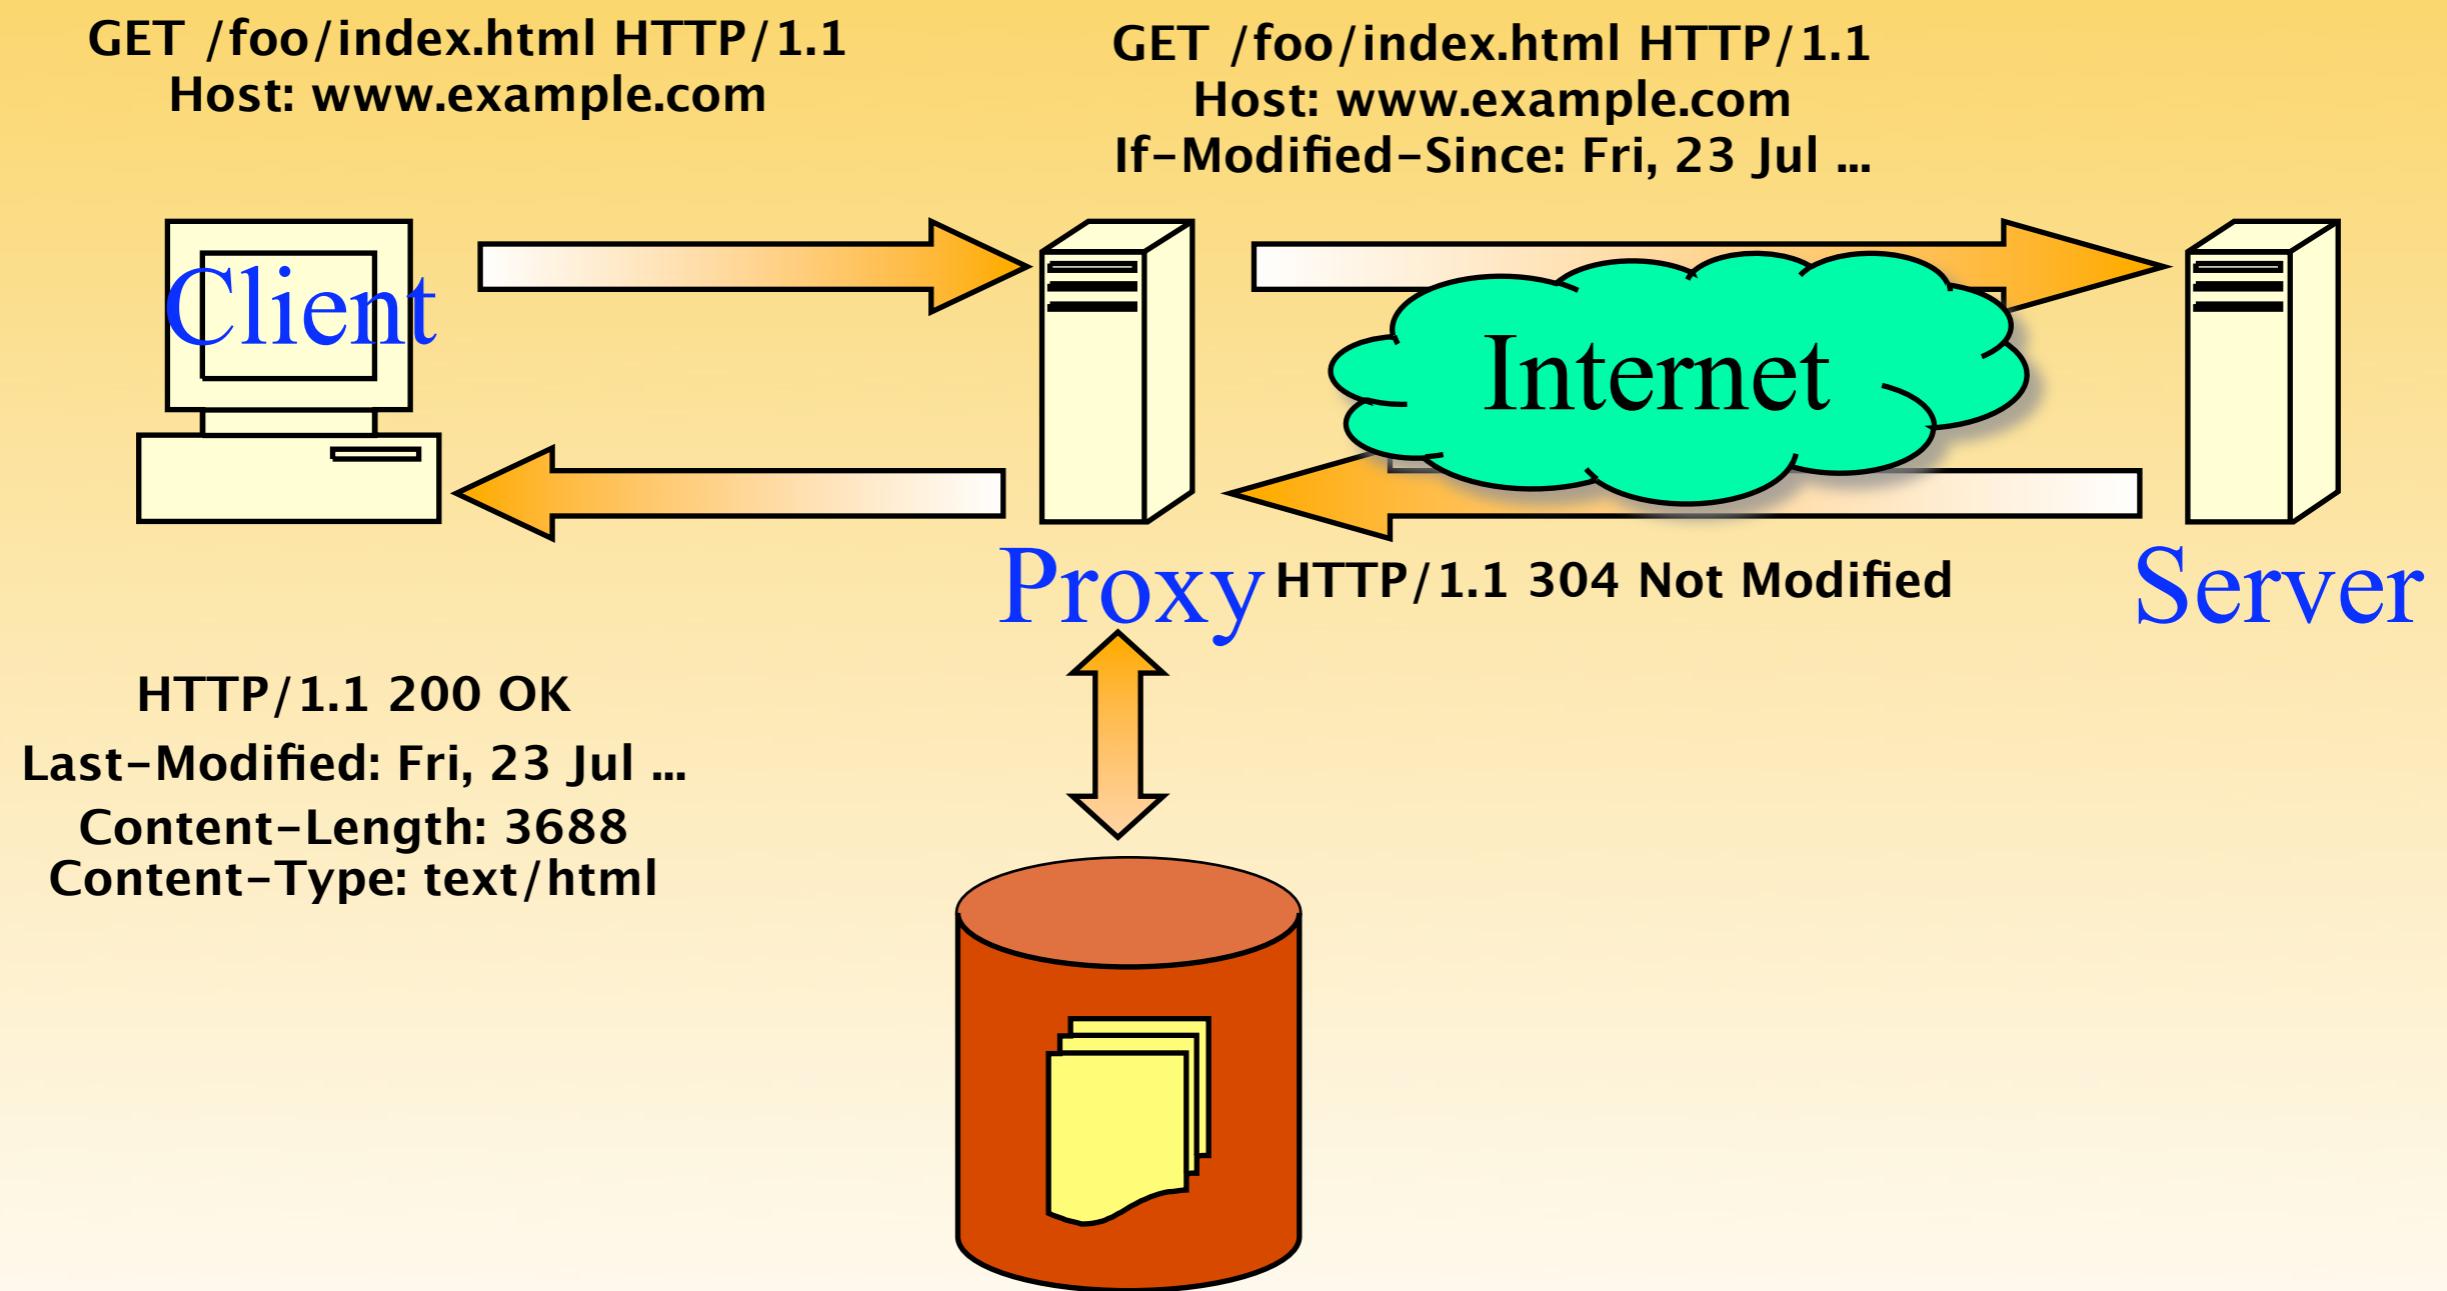


Proxy Cache Hit

GET /foo/index.html HTTP/1.1
Host: www.example.com



Proxy Cache Revalidation



Expires HTTP Header

- Tells all caches how long the associated representation is fresh for
- Set on web server
 - Absolute time (based on either last access or last modification time)
- Good for static images
- Also good for dynamic content
- Must be an HTTP Date; otherwise considered past
- Important for clocks on cache and server to be sync'd

Cache-Control HTTP Headers (I)

- **max-age=[seconds]** — specifies the maximum amount of time that an representation will be considered fresh. Similar to Expires, this directive is relative to the time of the request, rather than absolute. [seconds] is the number of seconds from the time of the request you wish the representation to be fresh for.
- **s-maxage=[seconds]** — similar to max-age, except that it only applies to shared (e.g., proxy) caches.
- **public** — marks authenticated responses as cacheable; normally, if HTTP authentication is required, responses are automatically uncacheable (for shared caches).

Cache-Control HTTP Headers (2)

- **no-cache** — forces caches to submit the request to the origin server for validation before releasing a cached copy, every time. This is useful to assure that authentication is respected (in combination with public), or to maintain rigid freshness, without sacrificing all of the benefits of caching.
- **no-store** — instructs caches not to keep a copy of the representation under any conditions.

Cache-Control HTTP Headers (3)

- **must-revalidate** — tells caches that they must obey any freshness information you give them about a representation. HTTP allows caches to serve stale representations under special conditions; by specifying this header, you're telling the cache that you want it to strictly follow your rules.
- **proxy-revalidate** — similar to must-revalidate, except that it only applies to proxy caches. For example:
 - Cache-Control: max-age=3600, must-revalidate

Assumptions about content types

Rate of change once published

Frequently Occasionally Rarely/Never

HTML

CSS

JavaScript

Images
Flash
PDF

Dynamic Content

Personalized

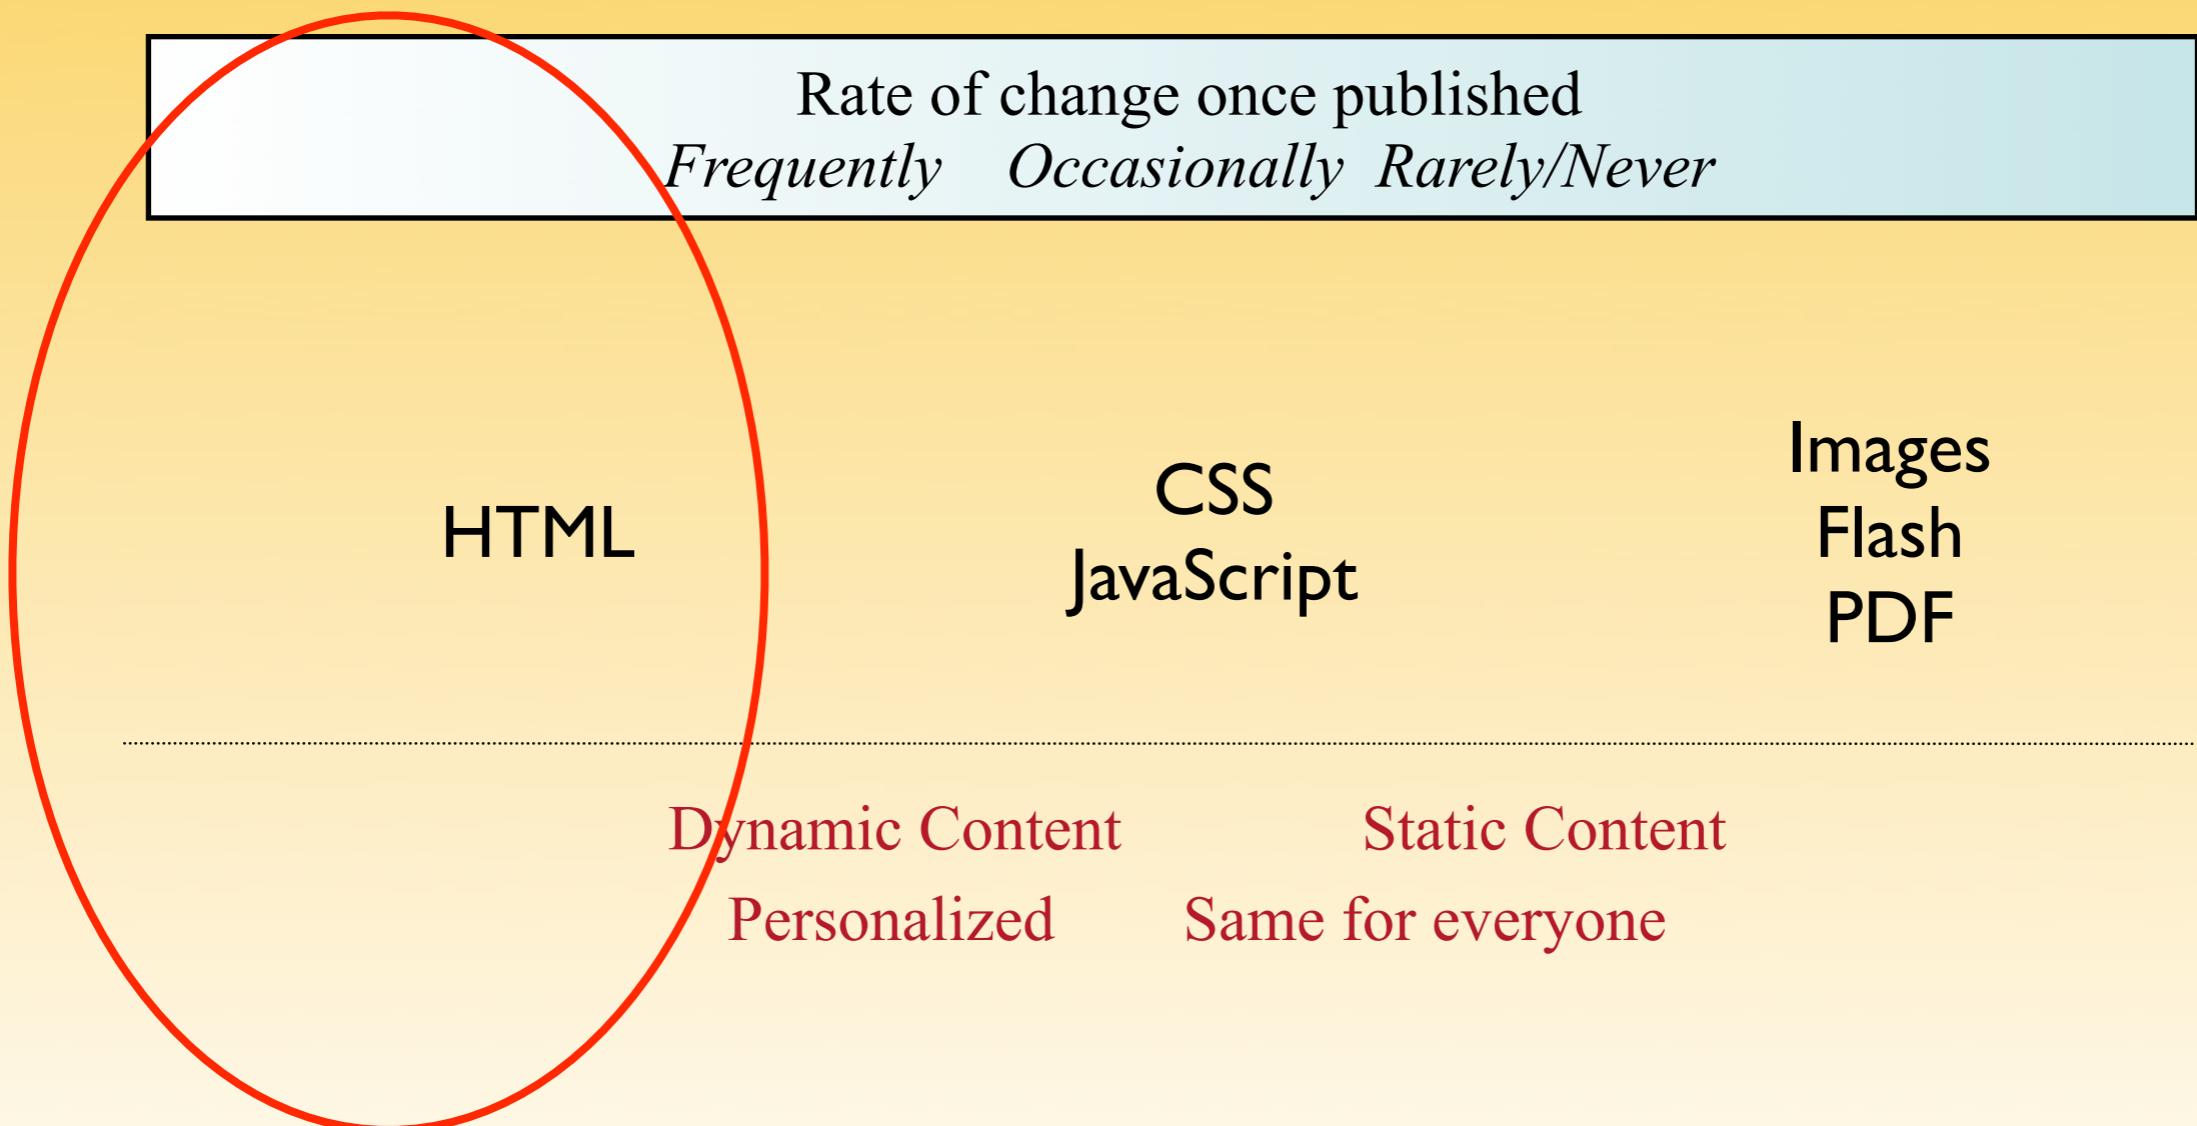
Static Content

Same for everyone

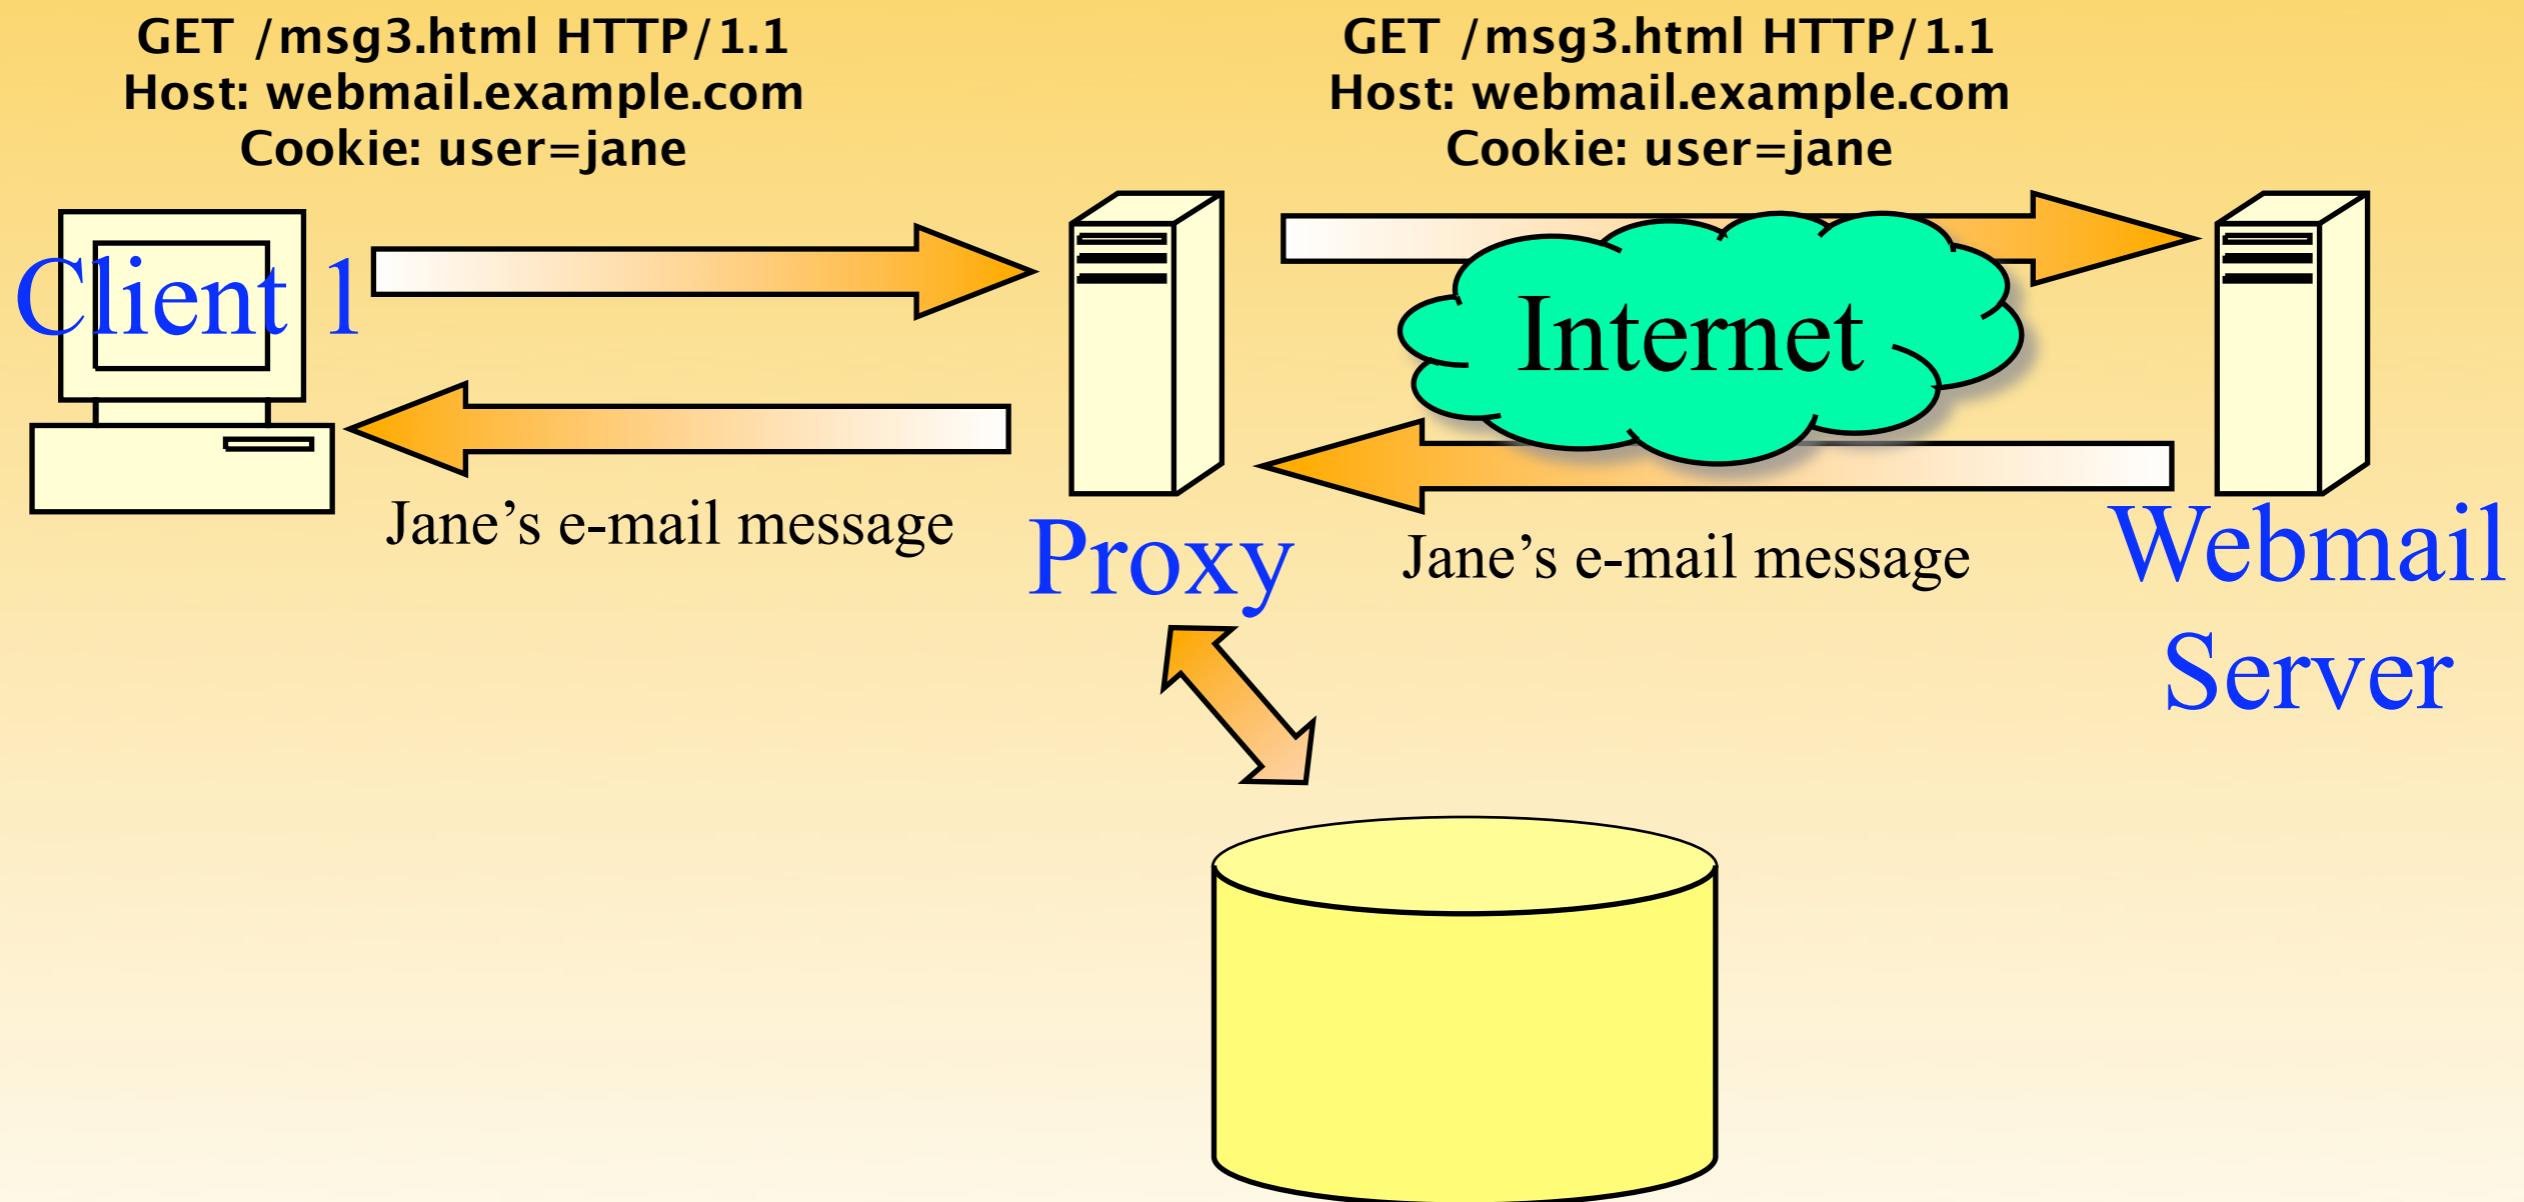
Top 5 techniques for publishers

1. Use “Cache-Control: private” for personalized content
2. Implement “Images Never Expire” policy
3. Use a cookie-free TLD for static content
4. Use Apache defaults for CSS & JavaScript
5. Use random strings in URL for accurate hit metering or very sensitive content

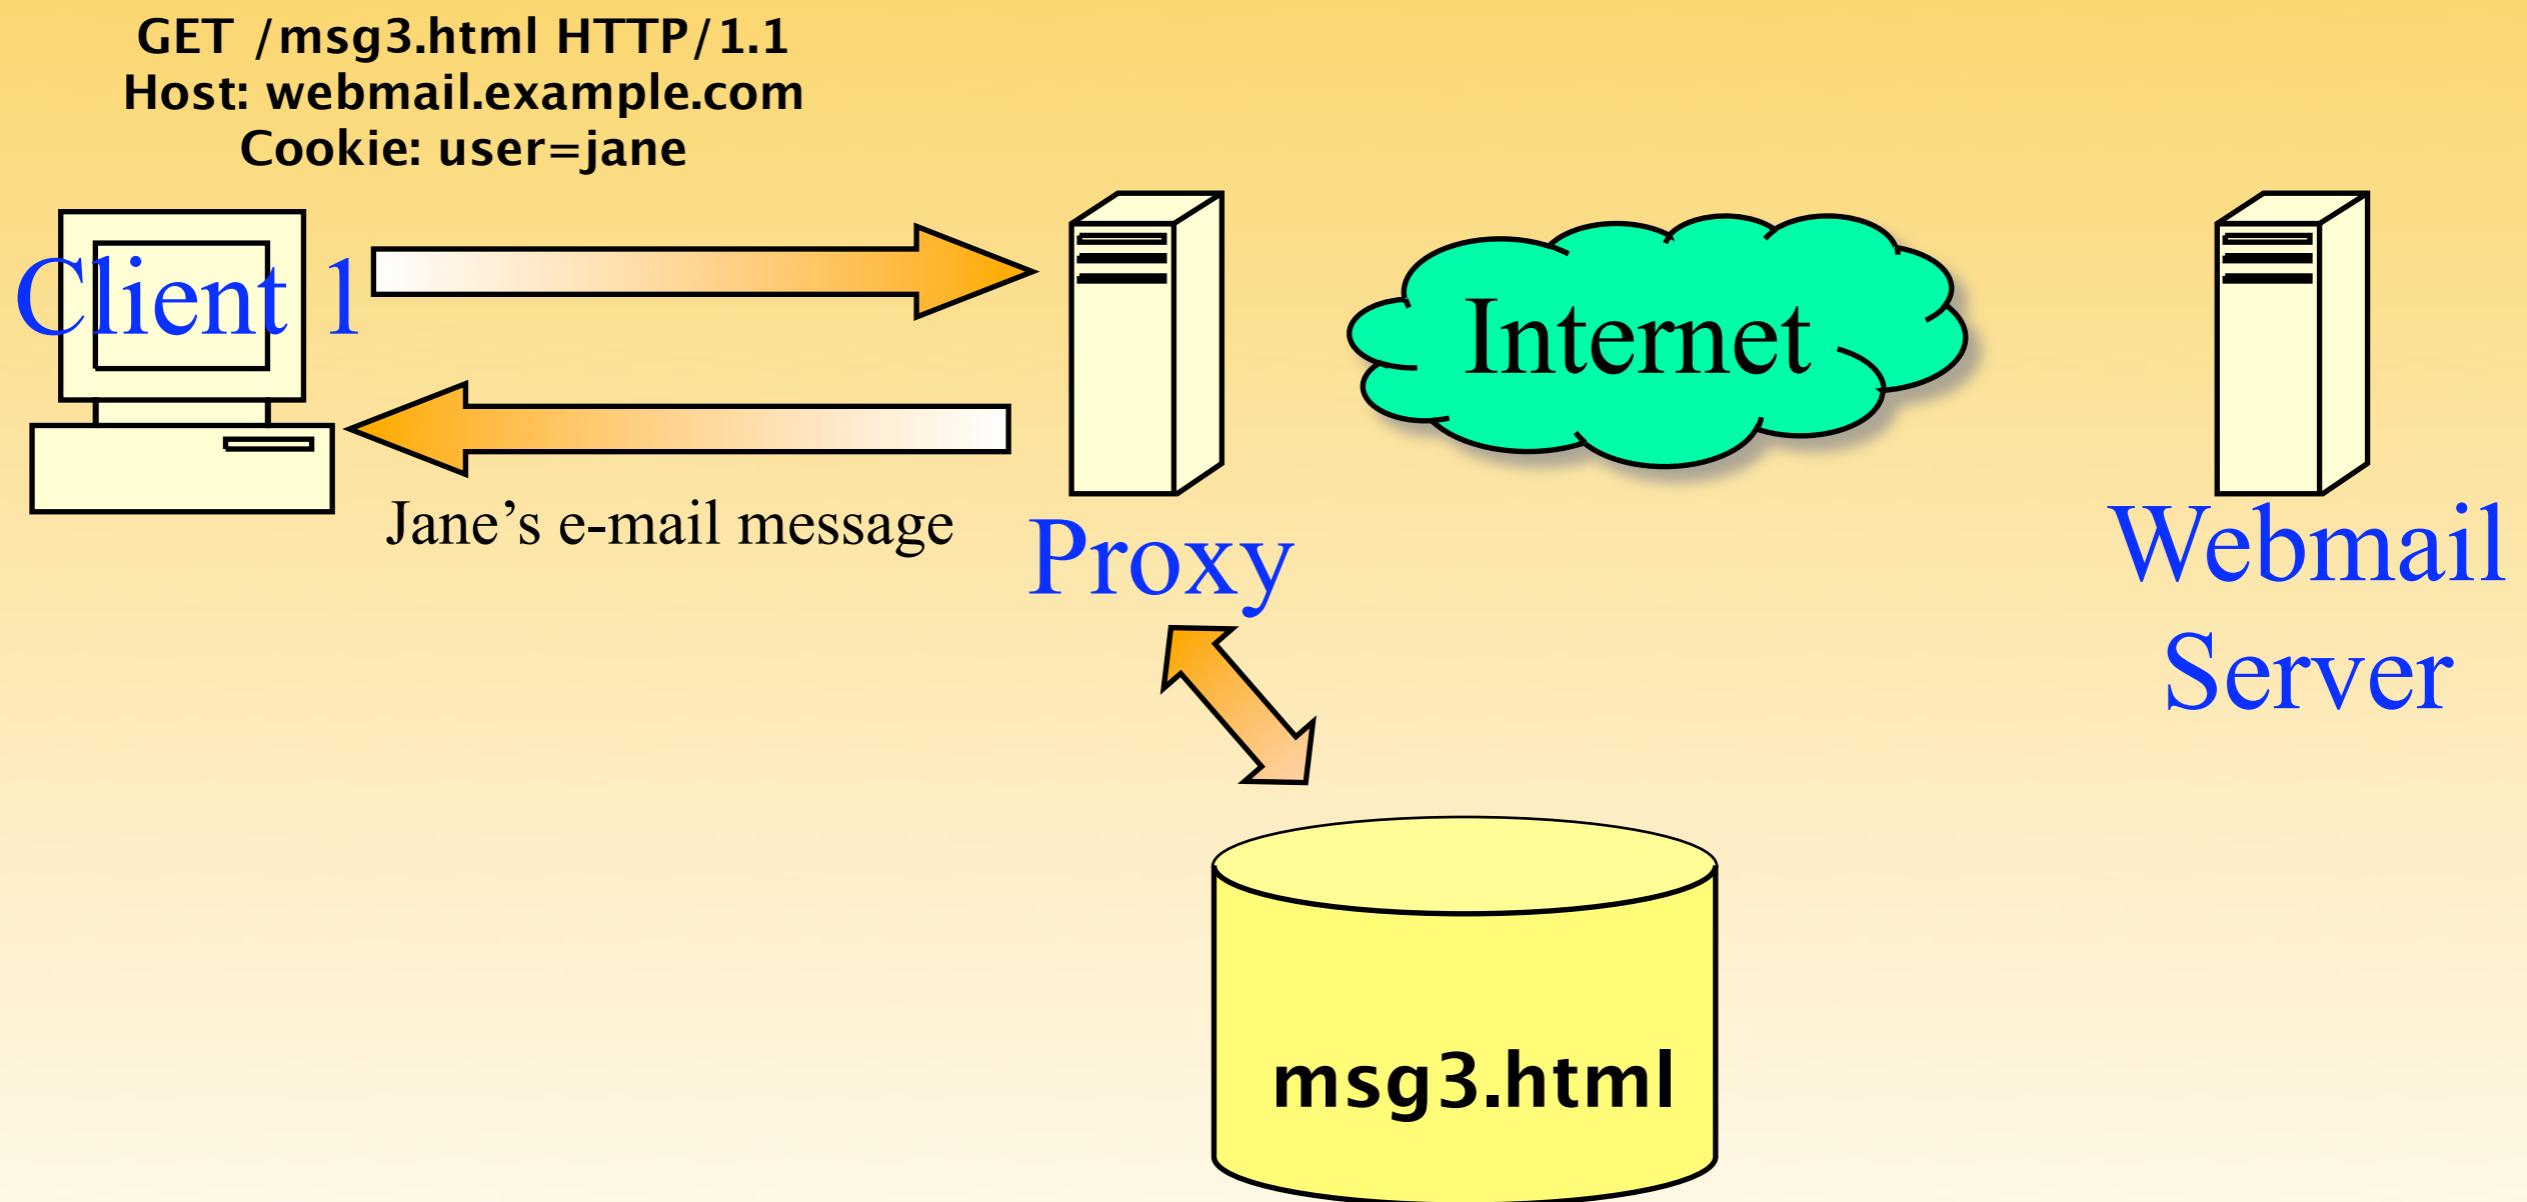
I. Use “Cache-Control:



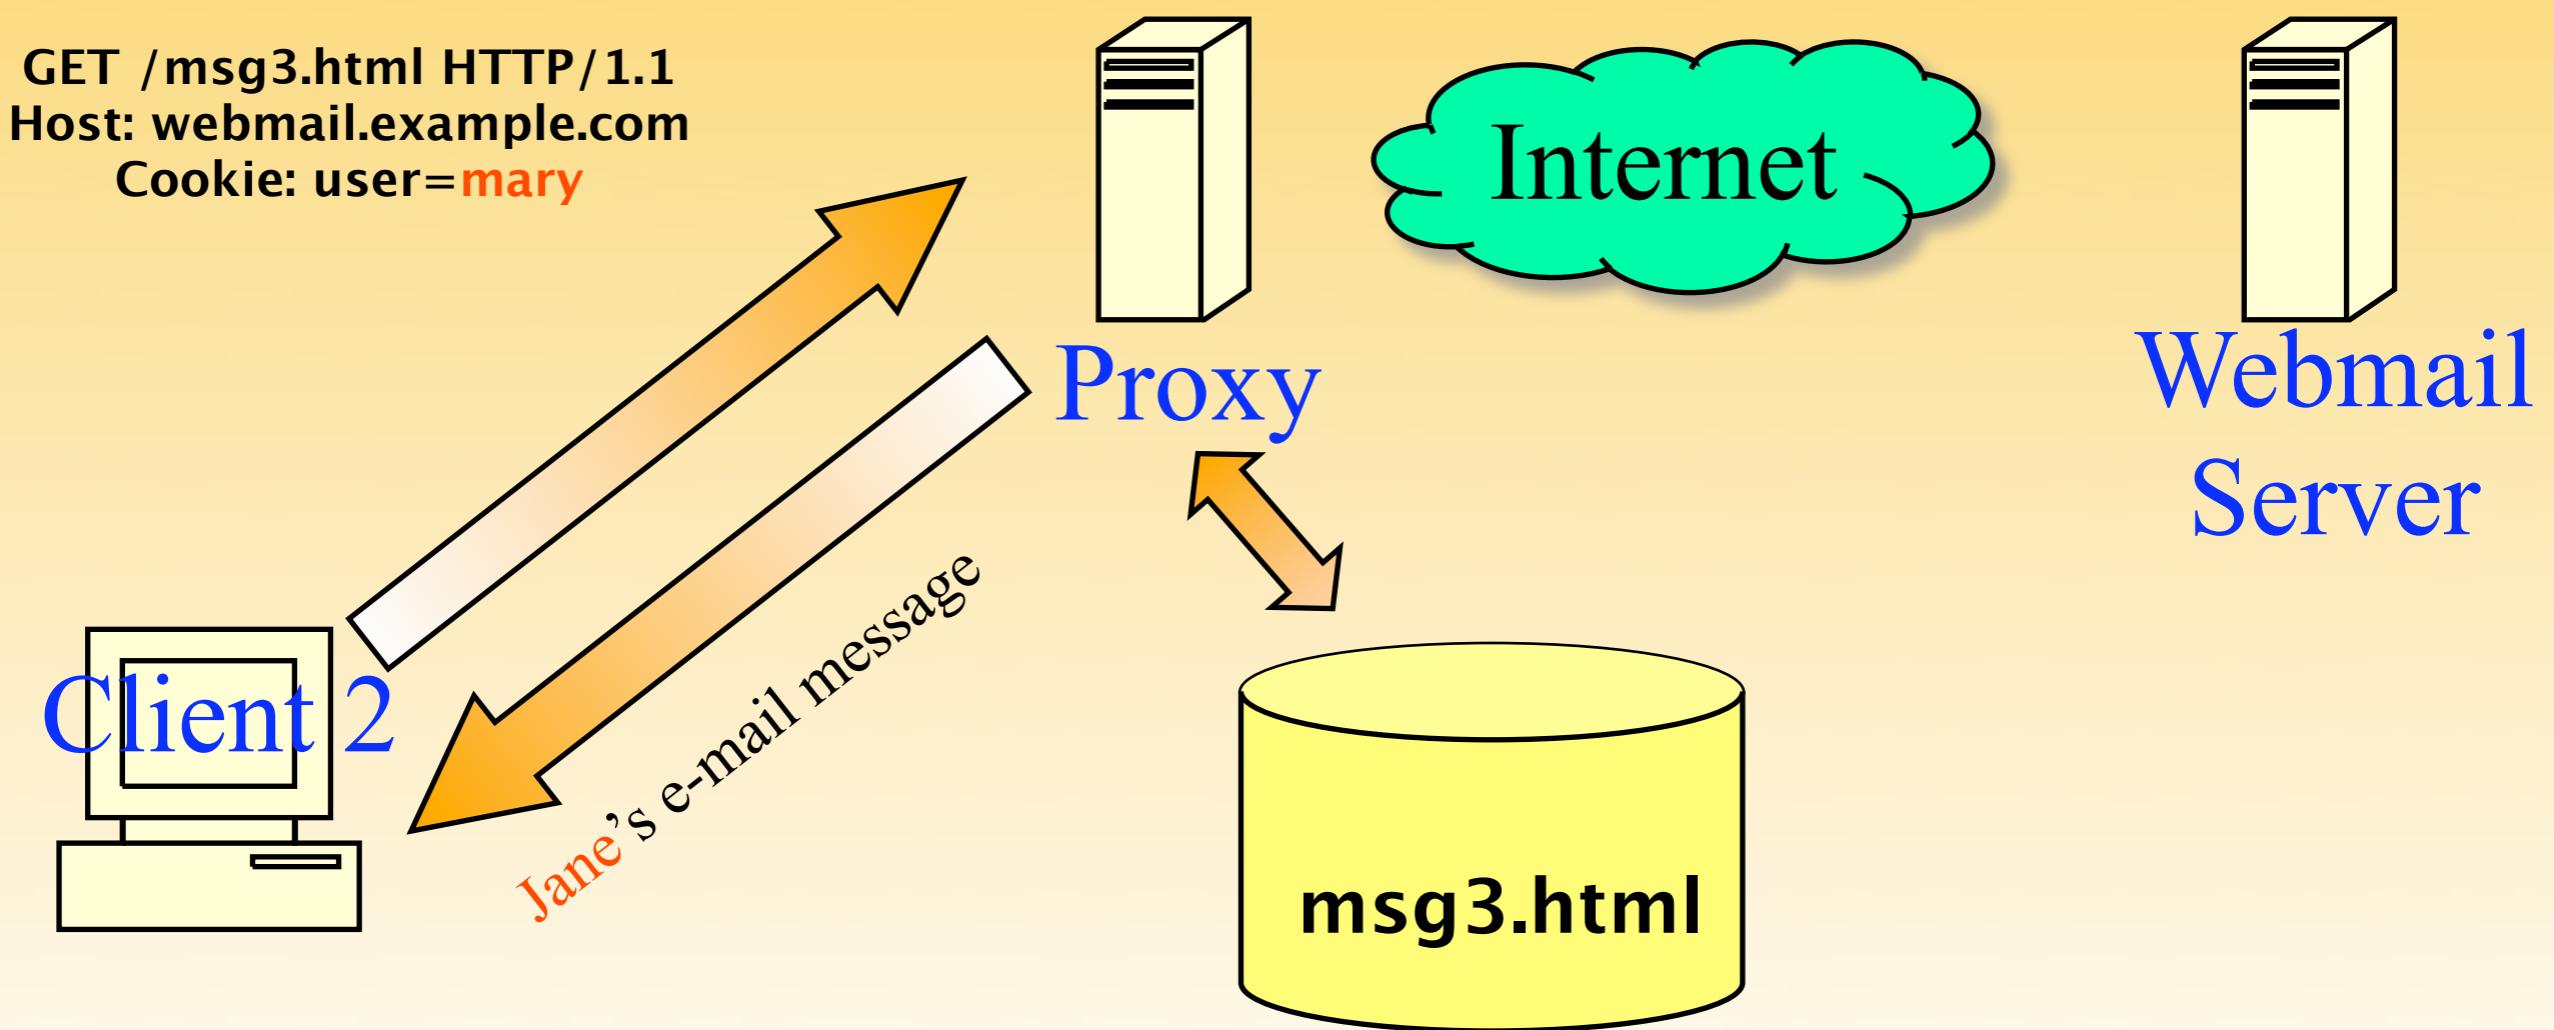
Shared caching gone awry



Shared caching gone awry



Shared caching gone awry



What's cacheable?

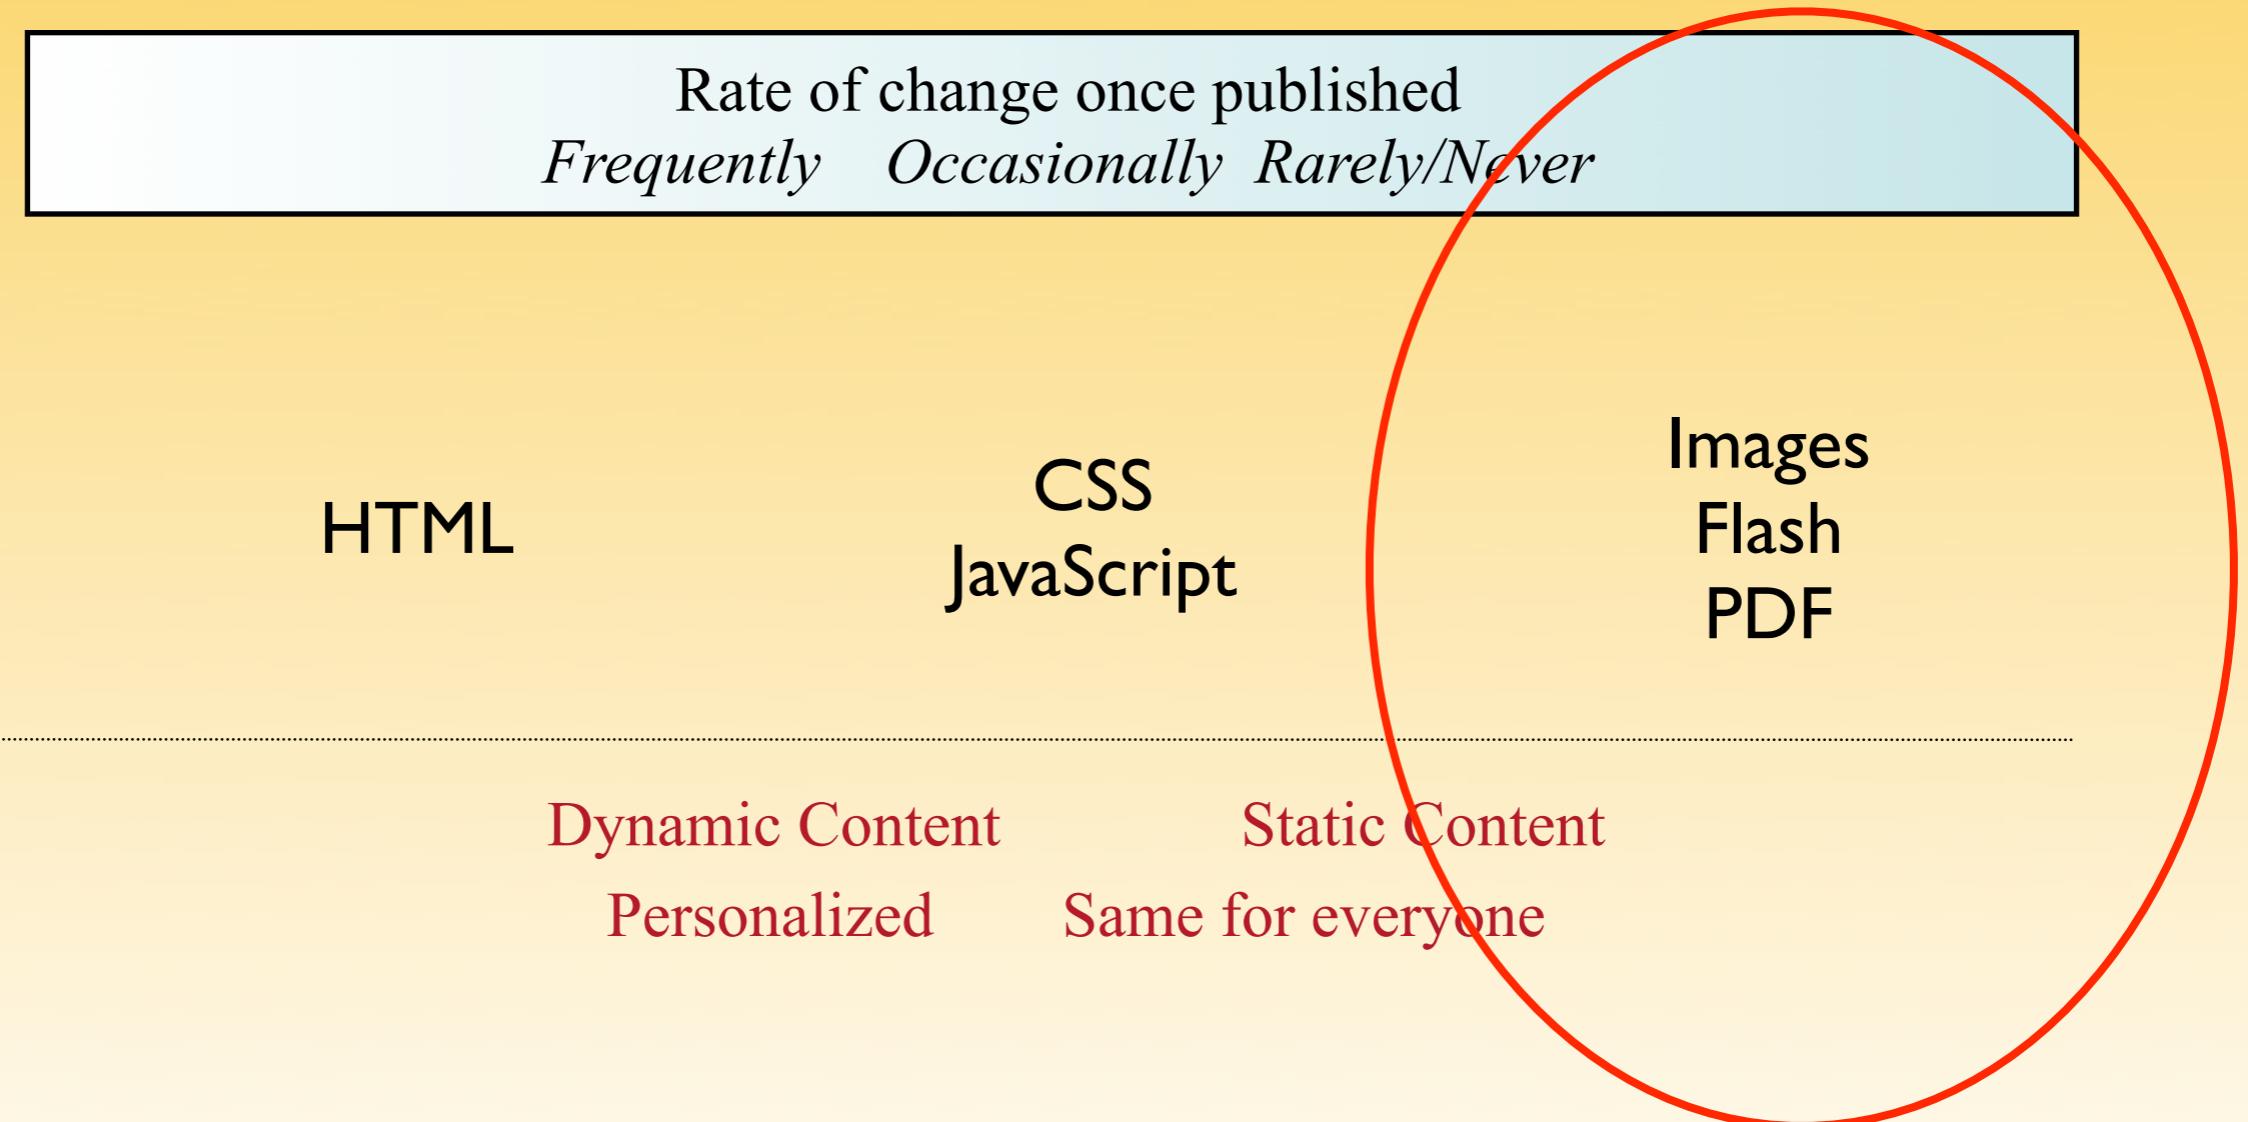
- HTTP/1.1 allows caching anything by default
 - Unless explicit Cache-Control header
- In practice, most caches avoid anything with
 - Cache-Control/Pragma header
 - Cookie/Set-Cookie headers
 - WWW-Authenticate/Authorization header
 - POST/PUT method
 - 302/307 status code

Cache-Control: private

- Shared caches bad for shared content
 - Mary shouldn't be able to read Jane's mail
- Private caches perfectly OK
 - Speed up web browsing experience
- Avoid personalization leakage with single line in httpd.conf or .htaccess

Header set Cache-Control private

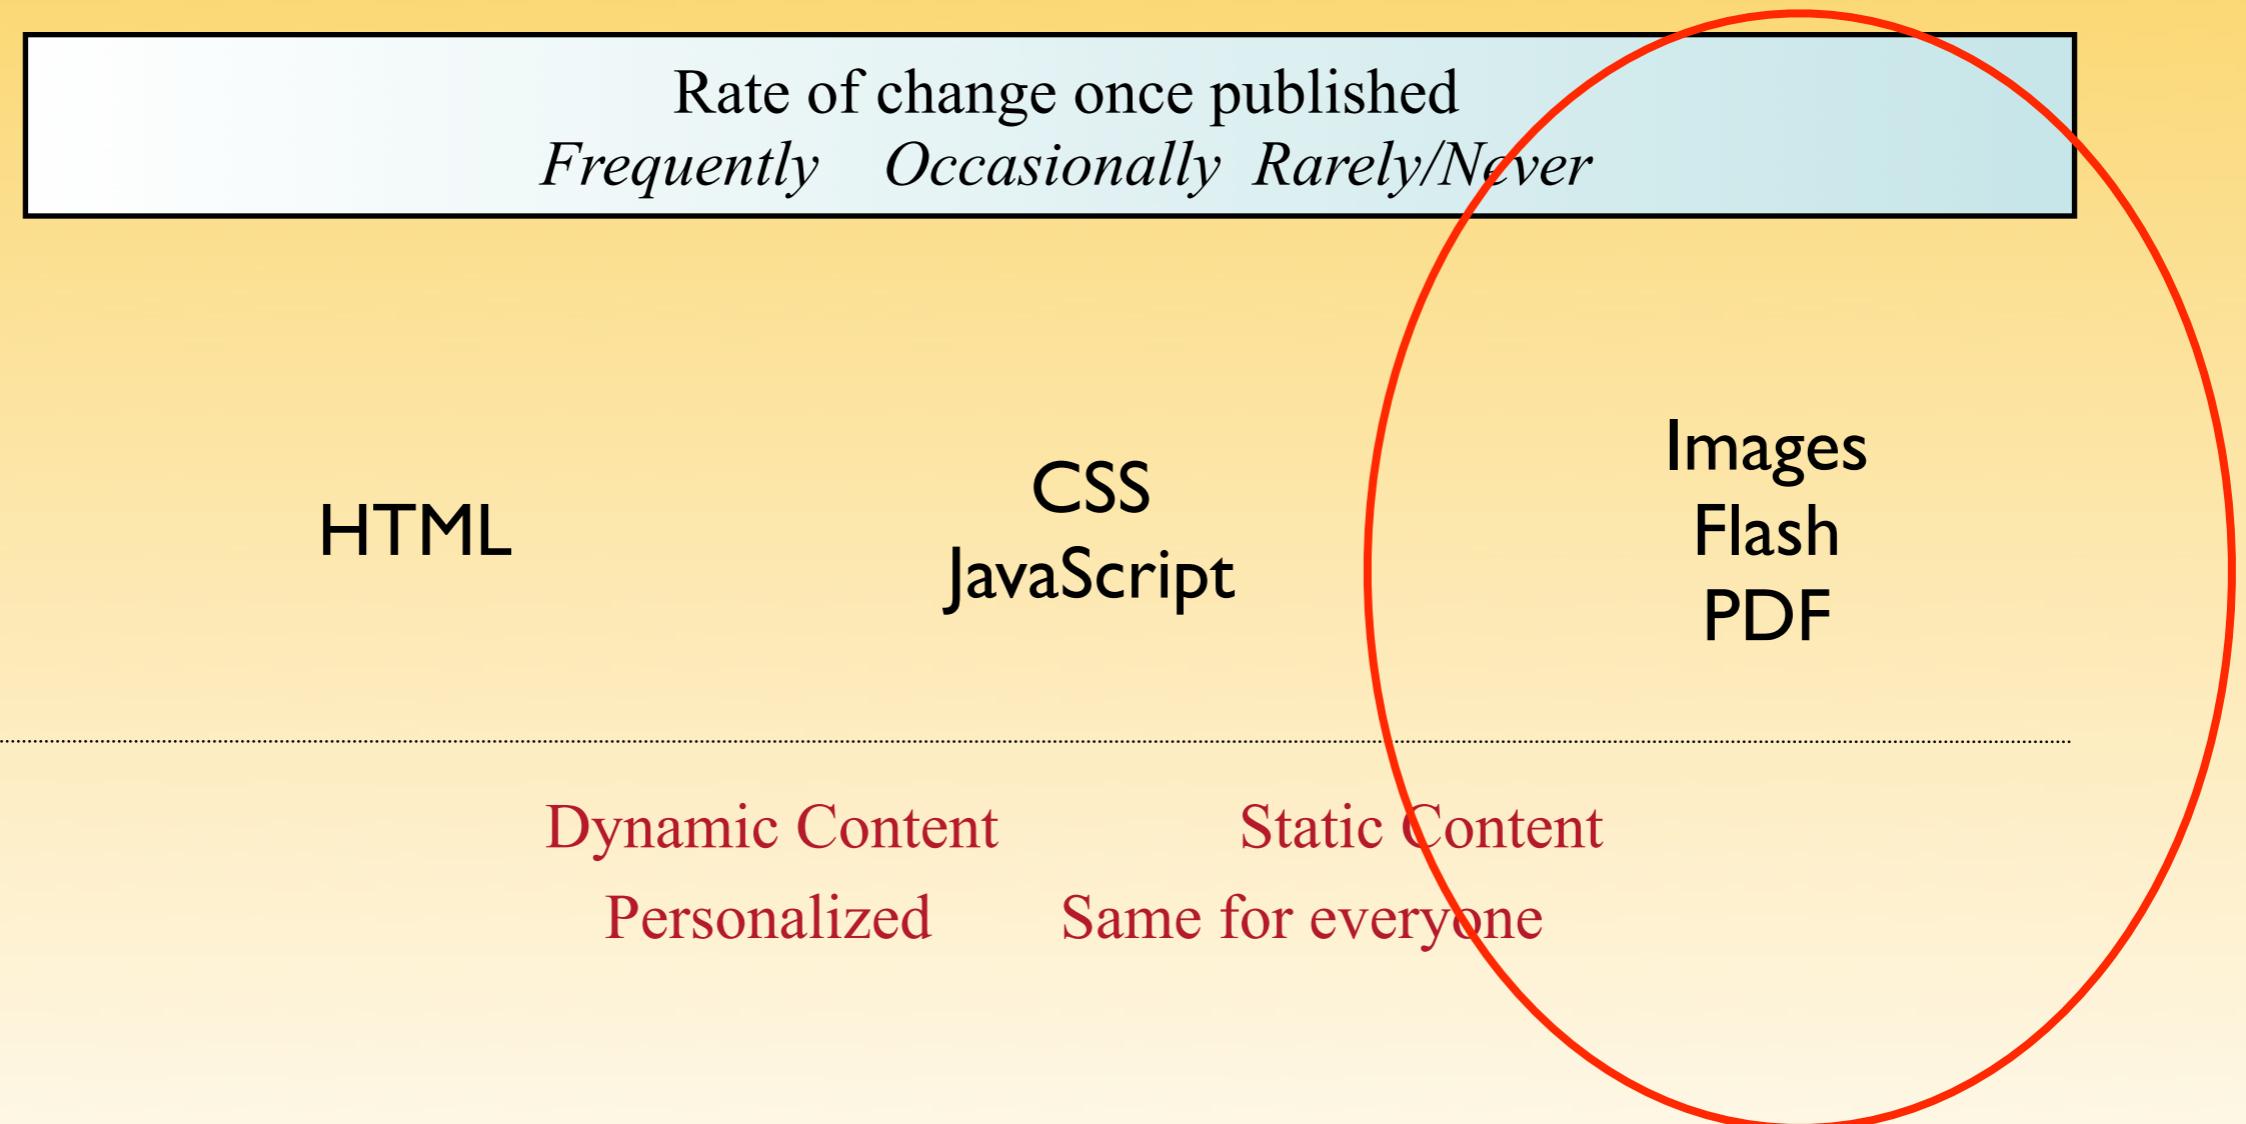
2. “Images Never Expire” policy



“Images Never Expire”

- Encourage caching of icons & logos
 - Forever ≈ 10 years in Internet biz
- Must change URL when you change img
 - <http://us.yimg.com/i/new.gif>
 - <http://us.yimg.com/i/new2.gif>
- Tradeoff
 - More difficult for designers
 - Bandwidth savings, faster user experience

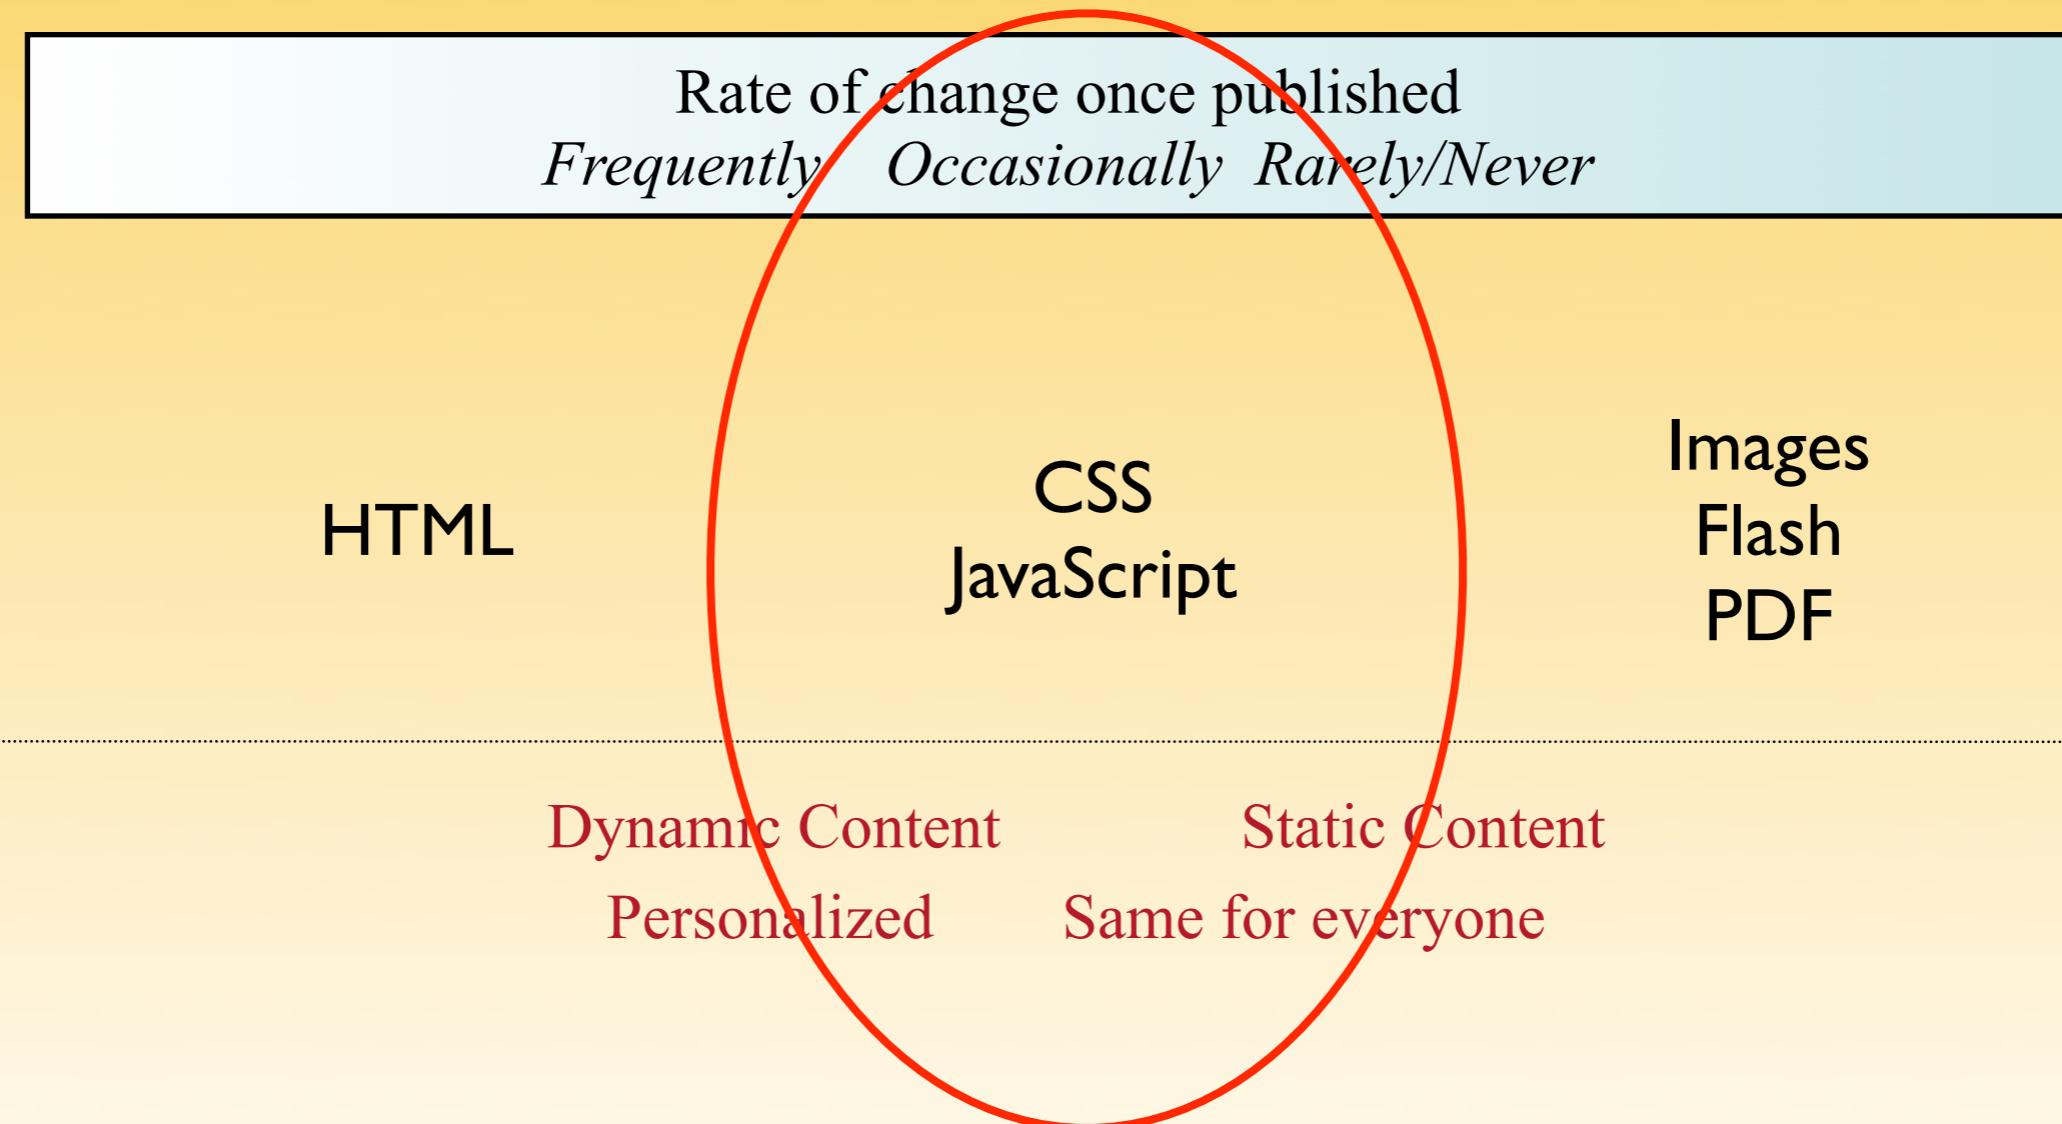
3. Cookie-free static



Use a cookie-free Top Level

- For maximum efficiency use two domains
 - `www.example.com` for HTML
 - `img.example.net` for images
- Some proxies won't cache Cookie reqs
 - But: multimedia is never personalized
 - Cookies irrelevant for images

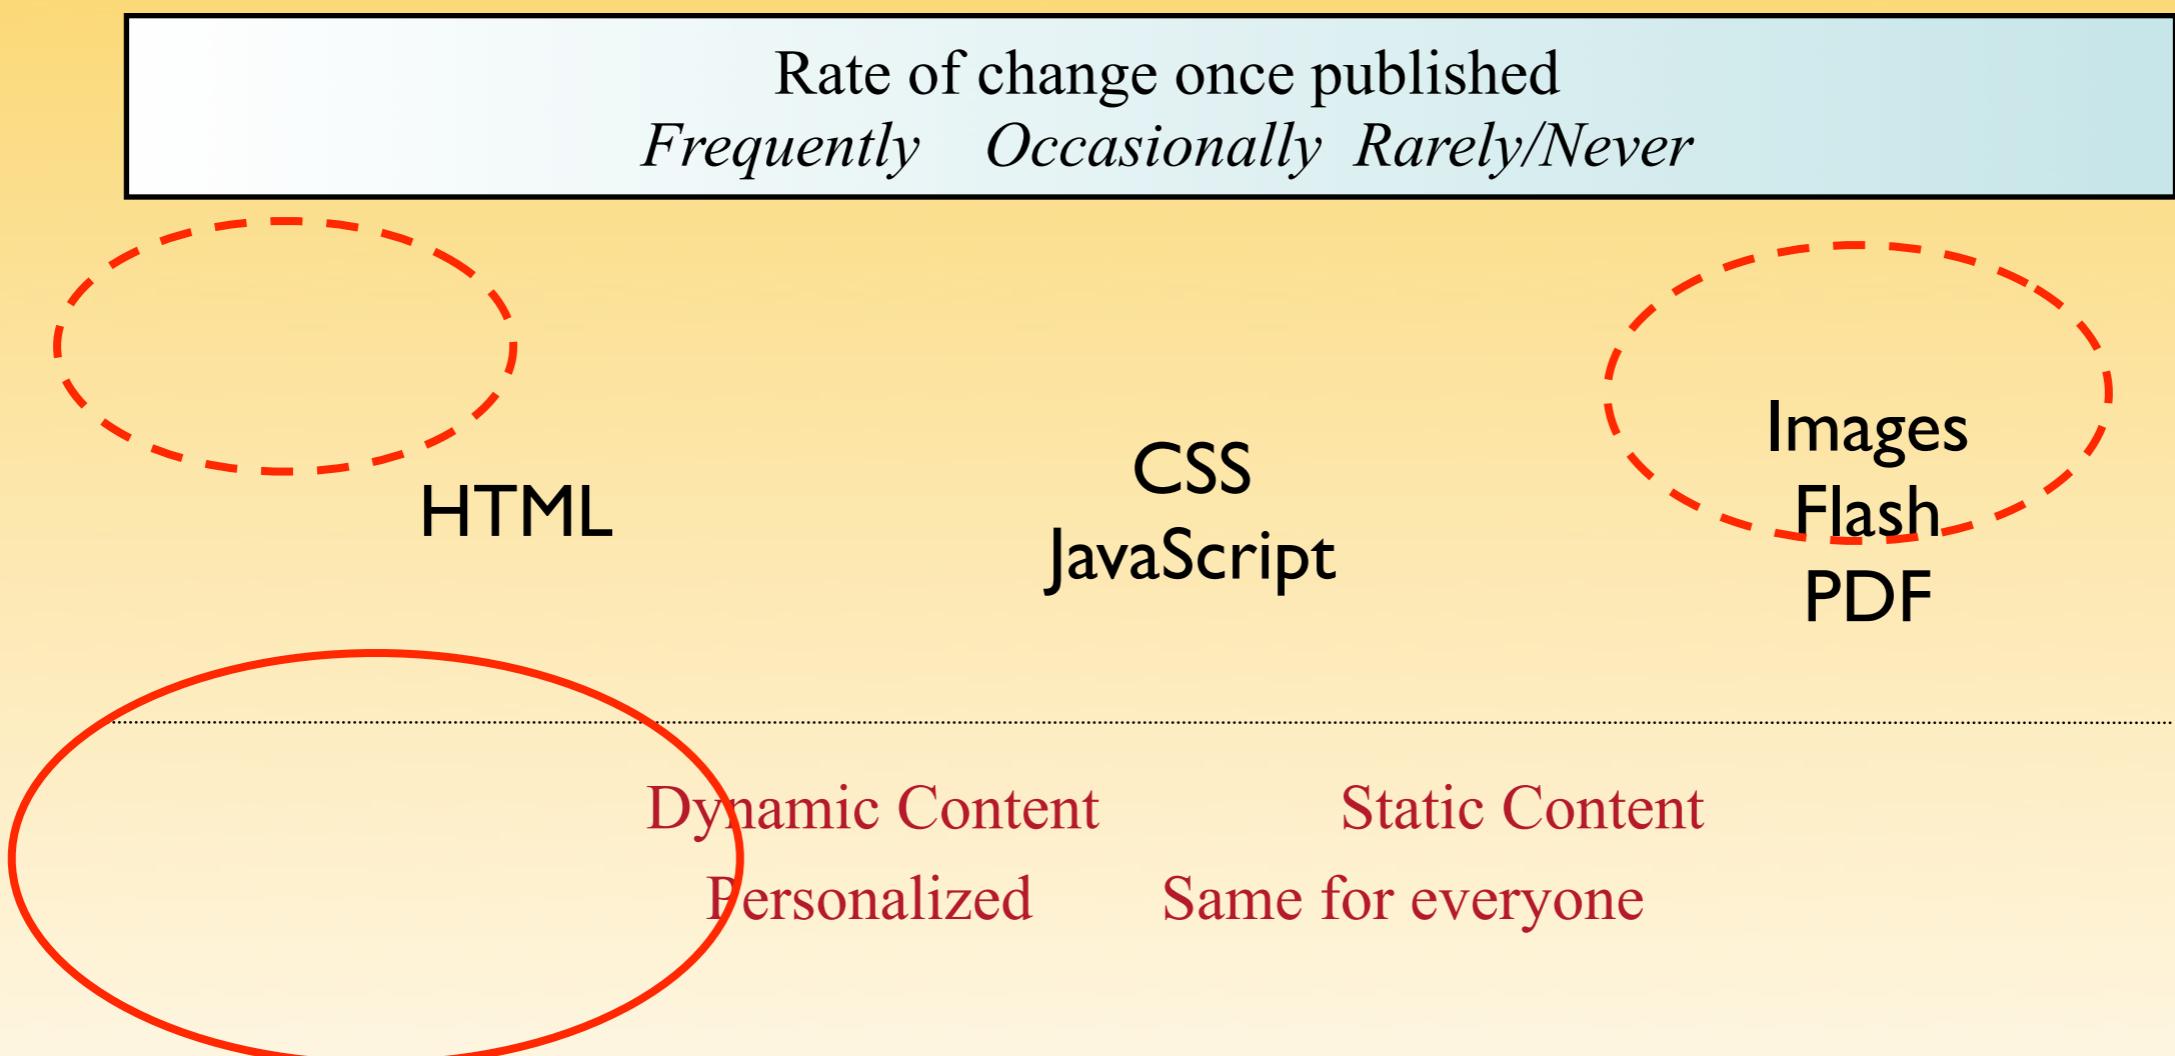
4. Apache defaults for



Revalidation works pretty

- Revalidation default behavior for static content
 - Browser sends If-Modified-Since request
 - Server replies with short 304 Not Modified
 - No fancy Apache config needed
- Use if you can't predict when content will change
 - Page designers can change immediately
 - No renaming necessary
- Cost: extra HTTP transaction for 304
 - Small with Keep-Alive, but large sites disable

5. Random URL strings for



Accurate advertising

- If you trust proxies
 - Send Cache-Control: must-revalidate
 - Count 304 Not Modified log entries as hits
- If you don't
 - Ask client to fetch uncacheable image URL
 - Return 302 to highly cacheable image file
 - Count 302s as hits
 - Don't bother to look at cacheable server log

Defeating proxies: turning public caches into private caches

- Use distinct tokens in URL
 - No two users use same token
 - Defeats shared proxy caches
 - Works well with private caches
- Doesn't break the back button
- May break visited-link highlighting
 - e.g. JavaScript timestamps/random numbers
 - Every link is blue, no purple

Breaking the Back button

- When users click browser Back button
 - Expect to go back one page instantly
 - Private cache enables this behavior
- Aggressive cache-busting breaks Back button
 - Server sends Pragma: no-cache or Expires in past
 - Browser must re-visit server to re-fetch page
 - Hitting network much slower than hitting disk
- Use very sparingly
 - Compromising user experience is A Bad Thing

Review: Top 5 techniques

1. Use “Cache-Control: private” for personalized content
2. Implement “Images Never Expire” policy
3. Use a cookie-free TLD for static content
4. Use Apache defaults for CSS & JavaScript
5. Use random strings in URL for accurate hit metering or very sensitive content

Review: encouraging caching

- Send explicit Cache-Control or Expires
- Generate “static content” headers
 - Last-Modified, ETag
 - Content-Length
- Avoid “cgi-bin”, “.cgi” or “?” in URLs
 - Some proxies (e.g. Squid) won’t cache
 - Workaround: use PATH_INFO instead

Review: discouraging caching

- Use POST instead of GET
- Use random strings and “?” char in URL
- Omit Content-Length & Last-Modified
- Send explicit headers on response
 - Breaks the back button
 - Only as a last resort

Cache-Control: max-age=0,no-cache,no-store

Expires: Thu, 01 Jan 1970 12:34:56 GMT

Pragma: no-cache